

Who Must Agree

Bond authorisation rules as coalition requirements

Empirical sections and exhibits — working conference draft

Compiled 25 August 2026 from the committed analysis record (corpus package v3, frozen).

*Draft 2026-08-24 (round 3). Every number traces to a committed RESULTS file (script-to-results map in **analysis/ANALYSIS_REVIEW.md**). [**PENDING**] marks claims held on a named data pass. The honesty record (nulls, pending items, data limits) is included below as Appendix H. House style applied throughout: impersonal constructions, no em dashes, British spelling. **The conference draft is FROZEN on corpus package v3** (the finance-flag fix, verified); package v4 and the B5 rerun land post-APSA only.*

Sections 1–3 (theory and institutional framework) are maintained separately and are not part of this compile. Reading paragraphs under each exhibit are drafting aids, marked strippable for submission.

4 · Data: observing the consent requirement end to end

Testing a theory of authorisation rules requires observing three things that no single dataset contains: the **rule** (who must agree, at what threshold), the **asking** (what governments put before voters, and what happened), and the **borrowing** (what was ultimately issued, under what security, for what purpose, and, critically, on whose authority). The analysis assembles all three at the level of the individual local government nationwide, and at the level of the individual referendum in nine states.

Referenda. The referendum panel compiles 47,235 local bond and tax measures from nine state administrative sources: California (CDIAC), Texas (Bond Review Board), Wisconsin (DPI), Louisiana (Secretary of State), North Carolina (NCSBE), Massachusetts (DLS Proposition 2 $\frac{1}{2}$ databank), Minnesota, Illinois and Indiana. Each measure is crosswalked to its Census of Governments unit: 40,924 (86.6%) resolve to a government identifier, with exact tiers verified at approximately 100% and the fuzzy-match tier independently audited at 95.1% accuracy. Restricting to measures at a genuine mandatory-ballot statutory cutoff leaves 23,577; the regression-discontinuity frame is its general-obligation subset, **11,889 GO measures** (Texas 8,062, California 2,189, Wisconsin 999, Louisiana 361, North Carolina 278) across three thresholds: 50% (TX, WI, LA, NC), 55% (California schools) and 66.7% (California non-school), with California thresholds assigned per measure from CDIAC.

The outcome architecture. The outcome of record throughout is the **official-statement corpus**: 258,762 official statements (2005–2025, all fifty states; 43,030 issuers), from which are extracted, for every issue, the par amount, the security pledge class (GO, revenue, lease, special tax), use-of-proceeds lines classified into 118 functional activities, and the variable the paper turns on, the ****authorisation mode evidenced in the document itself****: voter approval (with the election date the document cites), governing board action, or statutory authority. It is the richest available measure of what was borrowed, for what, and on whose authority, and every headline provision claim in the paper (issuance, channel, purpose, project continuation) rides on it. The Census survey data enter in two deliberately subordinated roles, each occupying a single table row: the Government Finance Database debt items (1967–2023; 2.1 million local unit-years, extended to fiscal 2024 by the public-use continuation) serve one purpose, a survivorship robustness check on the disclosure corpus; and the survey capital-outlay items serve one purpose, a bound on the channel the corpus cannot see by construction, capital spending without debt (the pay-go bound, Section 6). Neither survey use displaces the corpus as the outcome of record. The conference draft is pinned to corpus package v3 (verified 2026-08-24); later package versions enter only in the journal revision.

Validation. Because the authorisation mode is extracted rather than administrative, it is validated against the election record before use (Table 4.1). Where an official statement cites an election date, that date matches an independently observed referendum for the same unit 67.9% of the time pooled, and 95.4% in Wisconsin, the state whose registry is known-complete; conditional on a match, 91.3% of cited elections are passed measures (98.0% in California, 98.8% in North Carolina). Shortfalls concentrate exactly where registry coverage is known-short (Minnesota’s registry begins in 2020; Massachusetts dates votes at month grain), and the voter-mode support rate (76.2% of voter-mode new-money documents have a passed referendum within six years in the records) is a lower bound driven by authorisations older than the observation windows: Texas districts issue against decade-old voter authorisations in series. The survey side is validated in the other direction: the 2022 public-use file’s issuance item matches the GFD within 0.5% for

99.9% of bridged units.

The rules panel. State by entity-type by purpose by year coding of referendum requirements (threshold, ballot mandate, stringency). [PENDING] The panel is a first-pass machine coding cross-validated at 78% (29/37 codable states) against an independent hand coding; a 21-cell human verification pass is with the owner. Until it lands, every rule coefficient below is presented as a first stage (the association between the coded rule and observed authorisation behaviour), never as a causal estimate. The convention is marked in each table note.

Table 4.1 Validation of extracted authorisation fields. Table 4.2 Frame construction (universe, cross-walk, cutoff sample, GO RD frame), with per-state counts.

5 · The institutional landscape and the fifty-state first stage

Before estimating anything, the object itself is documented; the coded rules are then confronted with observed behaviour across the whole local state.

5.1 Three facts about the landscape

****Fact 1:** entity types hold radically different menus of exits from the voted channel.** Table 5.1 tabulates, for the national new-money corpus (78,672 canonical issues, package v3), the share of determined dollars authorised without a vote, by entity type. School districts route only **36.6%** of new-money dollars through non-voted channels; townships 62.0%; special districts 73.8%; counties **83.1%**; municipalities **88.7%**; and authority-class issuers (housing, health-hospital and utility conduits) **97.6%**. The full matrix (security class by authorisation mode, per entity and rule regime) shows why: a school district's menu is essentially GO-or-nothing, while a city can finance most of what it does through revenue bonds, leases and conduit authorities that never face an electorate. "Who must agree" is not one institution: it is a menu that varies by the kind of government doing the asking.

****Fact 2:** the balloted local state is the non-chargeable civic core; the chargeable perimeter never votes.** Classifying ballot purposes across 19,600 bond measures, K-12 schools alone account for 38.2% of everything put before voters, followed by water and sewer, roads, parks and fire protection (Table 5.2, Panel A). Panel B lists the corpus functions whose dollars are voted on less than 2% of the time: public hospitals (\$197.7B of local project dollars), multifamily housing (\$70.0B), electric generation (\$67.1B), gas utilities, airport terminals. A \$0.84B Harris County Hospital District line, San Francisco's multifamily housing revenue programme, Energy Northwest's generation debt: none faced a ballot; Los Angeles Unified's \$9.0B ask in 2024 did. The nationwide sorting is stark at the channel level: the voted channel carries **11.3%** chargeable dollars, the board channel 60.0%, the statutory channel 72.5%; measured in classified line counts, which are robust to the regime imbalance in line-level dollar coverage (appendix Table A-C1), the gradient is 17.9%, 34.5% and 54.0%. Composition exhibits report both bases. What the consent requirement governs, overwhelmingly, is the class of goods that cannot be charged to users, precisely the goods H2 says it should.

****Fact 3:** the same legal sentence assembles coalitions that differ by five orders of magnitude.** Where vote counts exist (TX, WI, LA, NC; 7,473 GO measures, since California's CDIAC reports percentages only), the median school bond was decided by **768** yes-votes, the median off-cycle city measure by **1,179**, and the median special-district election by **34**. Under the identical Texas

50% rule, Harris County’s road programme required 511,375 yes-votes of 1,022,748 cast, while a developer municipal utility district was authorised by **two votes of two cast** (Table 5.3). The coalition a referendum requirement demands is not a constant of the law but a variable of the electorate the law convenes, the observation that motivates the heterogeneity analysis of Section 8. (Texas rows carrying the registry’s counts-unknown placeholder convention, 3,188 rows recorded "1-0", are excluded and documented; no RD estimate is affected.)

5.2 The rule and the channel: the fifty-state first stage

A panel of **90,604 local governments** (every county, municipality, township, school district and special district in the national GFD universe) carries the coded rule for its class, demographic and fiscal characteristics (population or enrolment; revenue, property tax, debt stock; ACS homeownership, age, diversity and income at place grain for cities and county grain otherwise; county presidential partisanship), and its borrowing 2005–25 from the corpus and 2005–23 from the survey. A subpanel of roughly 570 large cities adds form-of-government institutions, tax-and-expenditure-limit stringency, city-level partisanship and mayoral party. Estimation: weighted least squares with state-clustered standard errors, Census-region fixed effects, and the characteristics above as controls.

The rule shows up in behaviour, most where the theory says it must. Under a strict rule, the OS-evidenced voted share of determined new-money dollars is higher by **+0.69 for school districts ($t = 4.5$)**, +0.20 for municipalities ($t = 5.8$), +0.18 for special districts ($t = 3.5$) and +0.11 (n.s.) for counties (Table 5.4); the pooled coefficient is composition-sensitive across classes and the per-class rows are the cited object. Raw magnitudes tell the same story: 68.4% of school dollars are voted under strict rules against 7.4% under lax ones; 22.2% against 3.7% for cities; 23.3% against 5.4% for counties. Figure 5.1 maps the resulting geography of consent across 46 states with full flag coverage: the observed voted share of local new-money debt ranges from roughly 80% in Oklahoma and two-thirds in Texas to under 3% in New York, Pennsylvania, Tennessee and Kentucky, the states whose statutes require no local ballot. The one reversal in the class table, townships (18.2% strict against 46.6% non-strict), is itself diagnostic: townships carry a proxy (municipality) rule in the panel, and New England towns borrow by town-meeting vote in states coded lax for cities. The panel is signalling that the coding needs a township class; it is flagged rather than absorbed. **[PENDING: rules pass-2; the township column is on its worklist.]**

Cities under strict rules substitute away from the voted instrument. At the unit grain with full controls, general-purpose governments in strict states carry a GO security share **27.7 points lower** ($t = -3.9$) than their counterparts in lax states; the non-chargeable project gap is **-8.7 points** ($t = -1.6$) in dollars and **-4.7 points ($t = -2.3$)** in classified line counts, the count basis being robust to the regime imbalance in classified dollar coverage documented in appendix Table A-C1 (Table 5.5 reports both; the text cites the count-based coefficient). This is the fifty-state generalisation of the sorting in Fact 2: where the GO instrument requires a coalition, governments that can finance through chargeable, unvoted instruments do. The state-level fixed-effects version of the same test gives -0.162 ($t = -1.83$). ****[PENDING: both are held as first-stage or descriptive until the rules pass-2.]****

The extensive margin is quiet. Whether a government borrows at all (2005–25) and how much it issues in survey totals are not significantly associated with the rule ($t \approx -1.3$). The rule moves the channel and the composition, not the existence of borrowing, exactly the pattern the

regression-discontinuity results of Section 6 rationalise.

Reform events: the honest state of play. The strongest available national causal design, California’s Proposition 39 (November 2000), which cut the school threshold from two-thirds to 55% while leaving every other California government and every other state untouched, does not yet deliver a verdict at the state-mean grain: the schools-only difference-in-differences is +0.45 log points (permutation $p = 0.15$ across 49 placebo states) with a visible pre-trend, and the sector-differenced triple-difference is approximately zero. Two data facts discipline the design: the GFD’s pre-2005 security split is degenerate for schools (full-faith-credit debt is 100% of school long-term debt by classification), so the certificates-of-participation-to-GO composition shift Proposition 39 should produce is invisible in Census data; and California’s state matching-bond waves (Propositions 47 and 55, 2002 and 2004) confound the totals margin. ****[PENDING: the publishable version is a district-level two-way fixed-effects design with enrolment weights, Conley-Taber inference, state matching-fund controls, and the CDIAC issuance database for the composition margin; a journal-version item.]**** The reform record accordingly plays its role in Section 9 as evidence on the politics of the rule, and the causal weight rests on the within-state designs of Section 6.

Table 5.1 *The menu matrix.* **Table 5.2** *The balloted and submerged local state.* **Table 5.3** *Absolute coalition sizes.* **Table 5.4** *First stage by entity class.* **Table 5.5** *Substitution.* **Figure 5.1** *The geography of consent (built: `analysis/fig_consent_map.svg`; entity-class small multiples in the appendix).*

6 · Authorisation at the margin: regression-discontinuity evidence

Design. For the 11,889 GO measures at statutory thresholds, governments that barely won authorisation are compared with those that barely lost it, in the threshold-centred vote margin. Estimation is local-linear with a triangular kernel; robust bias-corrected (RBC) inference is the lead convention (local-quadratic at h with robust variance), with conventional estimates alongside.

The naive gap, and what survives design discipline. Passed measures are followed by issuance 35.6% of the time against 14.6% for failed ones, a 21-point gap that is mostly selection. At the cutoff, the six-year effect of authorisation on GO issuance is ****+11.0 points (RBC, robust $z = 2.30$; conventional +14.6, $z = 4.47$)**** at the ± 10 bandwidth; any-issuance gives +11.6 and +14.4. The effect replicates state by state across three different thresholds: California +16.6 (z 3.0), Texas +14.1 (z 3.3), Wisconsin +21.2 (z 3.1). Three statutes, one answer. Louisiana’s negative cell (−27.6, z −1.9) is a data-grain artefact (measures fold to the parish); excluding it moves the pooled estimate to +16.2.

The intensive margin doubles, in two independent sources. Log new-money par per capita rises +0.92 (z 3.0) in the disclosure corpus and, decisively for the survivorship critique, **+0.83 (z 3.3) in the Census survey measure**, which includes the bank loans and private placements no official statement records; both are RBC-stable (+0.94, +0.88) and the survey pre-period placebo is null (+0.41, z 1.5). This is the survey debt items’ single appearance, as Section 4’s architecture prescribes.

Validity. Covariate balance: 0 of 7 pre-vote fiscal covariates imbalanced at ± 5 (max $|t| = 1.36$). Density: there is excess mass just above the cutoff (McCrary $\theta = +0.24$, $z = 4.6$); it is Texas-specific (TX +0.21, z 2.7; every other state approximately zero), carries the signature of endogenous proposal timing documented since Cellini, Ferreira and Rothstein (2010), and is itself agenda-margin evidence (Section 9). Identification survives it: donut estimates excluding the contested region are stable to larger (+0.19 to +0.20, z 3.7–5.3). Inference alternatives agree:

clustering by unit (z 4.11) or county (z 3.47); Lee bounds for crosswalk selection at trim 2.52% are $[+0.138, +0.163]$; randomisation inference in the ± 2 window rejects at $p < 0.0002$ (0 of 5,000 permutations). The two flexible-fit diagnostics that do not reject, the IK bandwidth ($h = 1.8pp$, $\tau \approx 0$) and RBC at ± 5 (confidence interval $[-0.09, +0.17]$), are reported with their reconciliation: on two points of support, slope terms absorb a level shift that the design-based tests at the same window detect decisively. The battery’s final row is **the pay-go bound**: rerunning the design with survey capital outlay per capita as the outcome, the schools post-minus-pre differenced estimate is **+0.377** (z 2.9; RBC +0.464, z 2.4) beside an issuance effect of +1.43, and trimmed levels imply an outlay-to-issuance ratio of **0.92** at the cutoff. Refused governments do not detectably replace lost borrowing with pay-go building (caveats: outlay spreads over construction years, so six-year windows understate long-project differences; the positive pre-period outlay RD from endogenous timing is netted by the differenced specification; total outlay only; full-window cohort = votes through roughly 2017–18, outlay ending at fiscal 2023). Nothing in the battery is hidden; the table shows all of it.

It is a timing effect, and the wedge is measurable. The event study loads the entire effect on the vote year ($\tau_0 = +0.236$, $z = 7.8$; every pre- and post-year approximately zero). Cumulative curves for the $|\text{margin}| \leq 5$ window (Figure 6.2) turn this into the paper’s quantities: among six-year issuers, the median barely-authorized government reaches the market in **0.33 years**, the median barely-refused one in **1.15**, a delay of **0.8 years**, and the refused side never reaches the authorized side’s end-of-vote-year issuance level within six years. By year six, 50.1% of barely-refused governments have still issued nothing, against 40.0% of barely-authorized ones: a ten-point wedge that is the durable residue after all catch-up.

Table 6.1 Main RD estimates. Table 6.2 The validity battery, including the pay-go-bound row. Figure 6.1 Event study. Figure 6.2 The cumulative wedge (built: `fig_cumulative_wedge.svg`).

7 · The response margin: what refusal buys

The theory’s distinctive claim is that a coalition requirement prices delay, not denial. Every refused measure is followed forward.

Refusal is a pause. Of 2,680 failed GO measures, the re-submission hazard is front-loaded (26.7% return within a year, then 22.8%, 15.8%, 12.4%), cumulating to **58.2% within four years**, with the median return at **1.02 years: the next election. Returns win: 61.9%** pass. Districts re-ask rather than concede: the median returning measure asks for **100%** of the original amount (n = 1,354; only 45.9% downsize).

The fate of the marginal refusal. For the 2005–19 cohort with full observation windows, per 100 barely-refused measures (Table 7.1): **54** are re-approved by voters within four years; 13 more return and wait; 5.2 borrow through the board or statutory channel; 9.0 issue on pre-existing voter authority; 18 are extinguished within the horizon. The 9.0 issuing on pre-existing voter authority draw on authorisations banked before the refused measure and are not evidence of evading it; the substitution reading attaches only to the 5.2 board-or-statutory cell. The transition matrix behind it shows what the cutoff actually moves: voter-mode first issuance (48.4% against 36.7%) against no issuance (43.1% against 54.2%), while board-mode first issuance is roughly 8% on both sides. The board channel is a floor, not the treatment margin. Consistent with this, the authorisation mix of what barely-refused governments still issue tilts towards the board channel (council share $\tau = -0.064$, $z = -2.0$).

The project itself survives. Matching ballot purposes to the use-of-proceeds functions of subsequent issues (deterministic category bridge, blind-audited at **80.0% precision and 88.9% recall**; every number cites the audit), same-purpose financing appears within six years for 44.7% of barely-passed measures and **33.4% of barely-failed ones** (RD +0.072, $z = 2.1$), with median arrival 0.32 against **1.67 years**. A third of narrowly refused projects get financed anyway, later. On re-submission, bundles are not recomposed: 78.2% of original purpose categories are retained. The pay-go bound (Section 6) closes the remaining loophole on the provision side: the same design run on survey capital outlay shows building falling with borrowing at a ratio near one, so the continuation gap measured in the corpus is not being quietly filled by debt-free construction. [PENDING: nine disputed audit pairs queued for the disagreement review.]

7.1 Authorisation banked, drawdown deferred

The supermajority regime produces a distinct fate for near-misses. Within California, failures with majority support short of the threshold re-submit far more than decisive failures (42.2% against 27.4%; 51.3% against 30.8% among schools) and show no jump at the symbolic 50% line: the response tracks proximity, not the majority label. Yet their issuance is lower, and the deficit is not a truncation artefact: it **widens** from -5.3 points at six years to -6.4 at eight (schools -6.7 to -8.0), under per-cell observability restrictions. The chain shows why. Near-miss failures convert by re-vote four times as often (107 against 25 passed returns), but the conversions sit undrawn: median pass-to-first-issue **6.2 years** against 2.9, and only 5.5% see a voter-mode issue within eight years. Where the bar is a supermajority, even re-assembled consent does not become borrowing on the study horizon.

A candidate mechanism, that Proposition 39's statutory tax-rate caps meter the drawdown of 55%-route authorisations, was tested by splitting Californian conversions by the route of the passing return: the uncapped two-thirds cell is too thin to test (9 school conversions against 118 capped) and its point direction runs against the cap story (median 9.7 years to first issue against 4.9 capped). The mechanism is neither supported nor excluded at this sample size; the fact stands documented and unexplained, with county assessed-value rolls and the CDIAIC sold-versus-authorised series as the upgrade path.

***Table 7.1** The fate table and transition matrix. **Table 7.2** Re-submission hazard, pass rates, amount ratios. **Table 7.3** Project continuation and bundle recomposition (with audit precision). **Table 7.4** The near-miss chain and the rate-cap split.*

8 · Where the requirement binds: exits and electorates

The RD average conceals the theory's structure. Two partitions recover it.

By exit menu. Splitting the RD frame by entity class and aligning each class against its independently measured national menu (Section 5, Fact 1): school districts (non-voted share 29.8%, window-chargeable menu 0.4%) show the binding effect (+0.147, z 4.0; RBC +0.123, z 2.3) and the highest re-submission rate (59.2%); general-purpose governments (non-voted share 80.4%) show no discontinuity at all (+0.073, z 0.9; RBC -0.02) and the lowest re-submission (25.3%); special districts sit between, with the largest GO-specific jump that shrinks at any-issuance as part of the gap reroutes to non-GO instruments. The special-district cells are small and the ordering, not the magnitude, is the cited object. The rule binds exactly where the menu

offers no exit; where exits exist, refusal is nearly costless and rarely re-litigated. This is the fork the fifty-state substitution results (Section 5.2) show in cross-section.

By electorate: the stable propertied public. The naive assembly-cost reading of H3 predicts that the authorisation effect weakens wherever coalition-building is harder, and in particular that demographic heterogeneity raises the price of assembly. The data falsify that version on its own sign: at the proper demographic grain, the effect is carried by homogeneous electorates (+0.232 against -0.018 in diverse ones), not blocked by them. What the moderator profile consistently favours is a **stable-propertied-public** account: authorisation binds where the consenting public is durable and propertied, because it is those electorates whose refusals stay refused and whose approvals licence long-horizon drawdown. Homeownership above the frame median gives $\tau = +0.182$ (**z 3.5**) against $+0.056$ (n.s.) below; the 65-plus share gives $+0.194$ against $+0.036$; and the effect is stronger on-cycle ($+0.252$ against $+0.099$ off-cycle), when the broad durable electorate is the one consulted. All moderation tables run on the 6,255 proper-grain measures (place grain for cities, district grain for schools); special districts are excluded pending sub-county demographics (TWDB interpolation), as each table notes. It is ownership, not affluence (the income split is flat: $+0.129$ against $+0.122$), and institutional, not partisan: precinct-built city partisanship and mayoral party in the 577-city panel show no moderation (all n.s.). The coarser county-grain partisanship splits (2020 vintage, county grain) agree and sit in the appendix. The reframing is stated as it is: an account adopted after seeing the homogeneity sign, offered with its transmission story, and flagged as post hoc rather than pre-registered.

The incidence of the higher bar. Blocked majorities (measures with majority support short of a supermajority threshold; California only, by construction) arise in places that sit between the decisive-failure and comfortable-passage places on renter share, diversity and inequality, and among schools they are the poorer places (median household income $-\$3.8k$ against cleared places, SE 1.5k). Descriptive, but it locates the supermajority's demographic cost: majorities in less affluent, mid-composition districts. ****[PENDING: within-matched ACS caveat as throughout; special districts remain on county-proxy demographics.]****

Table 8.1 The fork against the menu. Table 8.2 Moderator splits. Table 8.3 Demography of blocked majorities.

9 · The agenda margin and the politics of the rule

Rules discipline what is asked. Under California's 55% bar, school districts bring **8.8 proposals per 100 districts per year** at a median \$39M, 62.7% on-cycle; under Texas's 50% bar, **20.1** per 100 at \$15M, 14.1% on-cycle; Wisconsin sits between on all three margins. Fewer, larger, better timed: the coalition requirement operates before any vote is cast. Pass rates make the same point in reverse, nearly invariant between the 50% and 55% regimes (79.2% against 77.6%) and collapsing only at two-thirds (47.6%): a moderate bar is absorbed by the agenda; an extreme bar defeats even what selection permits. The Texas 2019 reform (HB 3) closes the loop within-state: when the state mandated separate propositions for stadiums, natatoria and performing-arts facilities, propositions per election jumped from 1.5–1.7 to 2.0–2.4 and the multi-proposition share doubled within two cycles. Bundling is a rule-governed choice, not a habit. The near-cutoff excess mass of Section 6 belongs to the same family: proposals are timed to win.

At fixed voter support, the rule alone moves outcomes. In the 50–55% band, the identical electoral result is a failure for a California school district and a success for a Texan or

Wisconsin one. The difference-in-differences across the adjacent band gives $^{**}+11.5$ points of six-year issuance (SE 4.8) ** attributable to the rule at fixed support, indistinguishable from the RD estimate obtained from an entirely different comparison: an out-of-design calibration regarded here as the strongest single check in the paper.

The polity fights over the rule exactly where the theory predicts. Since 1990, nearly every attempt to lower a local borrowing threshold has targeted schools, the class the menu matrix shows holding the poorest exit menu: California’s Propositions 170 (1993, failed), 26 (2000, failed) and 39 (2000, passed 53–47, schools only, with a \$30M elite-financed bipartisan campaign); Washington’s 2023–25 school-bond bills (all stalled); Idaho’s roughly eleven legislative attempts (none passed). The one recent non-school attempt, California’s Proposition 5 (2024), extending the 55% bar to housing and infrastructure, failed. Where exits exist, no one spends thirty million dollars to lower the bar. ** [PENDING: the reform table is compiled from secondary sources and marked `secondary_unverified` pending primary-record verification; the national event-study estimates for Proposition 39 remain inconclusive at the state grain (Section 5.2). The reform record is cited here as evidence on reform politics, not as a causal estimate.] **

***Table 9.1** Proposal behaviour by regime; pass-rate invariance. **Table 9.2** The institutional wedge. **Table 9.3** The reform record (appendix). **Figure 9.1** TX-2019 unbundling.*

Main tables

Table 1: Sample and summary statistics

State (threshold)	N	Mean yes	Passed	Median amount	Median electorate
CA (55%)	1,715	62.423	0.809	\$44.0m	–
CA (66.7%)	466	66.442	0.509	\$20.0m	–
LA (50%)	361	59.796	0.770	–	1,357
NC (50%)	278	65.843	0.914	–	30,568
TX (50%)	8,062	68.849	0.792	\$15.0m	1,397
WI (50%)	999	54.070	0.661	\$12.4m	2,988
All compiled referenda	47,235				
Matched to a Census government	40,924				
At a genuine mandatory-ballot requirement	23,577				
GO regression-discontinuity frame	11,889				

Panel B: the official-statement corpus (package v3)

Documents: 258,762

Issuers: 43,030

Determination rate: 0.937

Mode shares (docs): voter / board / statutory: 0.195 / 0.666 / 0.063

New-money par: schools / municipal / county / special / township (\$B): 524 / 606 / 220 / 139 / 35

Notes: Panel A: the referendum frame by state and threshold (cells ≥ 30 measures); amounts where registries print them (CA/TX/WI); electorate = votes cast where counts exist (TX/WI/LA/NC; counts-unknown placeholder rows excluded). Restriction cascade beneath. Panel B: corpus package v3.

Reading. The frame spans three statutory thresholds; the corpus supplies the authorisation mode for 93.7 per cent of a quarter-million documents.

Table 2: Covariate continuity at the authorisation threshold

Outcome	RBC	robust p	Conv.	h	Eff. N (L/R)
Any new-money issue, year $t-2$	+0.030 [-0.028, +0.089]	0.314	+0.007	± 10	1,596/2,743
Any new-money issue, year $t-1$	+0.021 [-0.036, +0.079]	0.468	+0.022	± 10	1,596/2,743
$\ln(1+\text{LTD issued p.c., } t-3..t-1)$	+0.530 [-0.256, +1.315]	0.186	+0.412	± 10	1,495/2,553
$\ln$ population	+0.320 [-0.521, +1.160]	0.456	+0.142	± 10	252/497
$\ln$ enrolment (schools)	-0.035 [-0.343, +0.272]	0.822	-0.001	± 10	1,285/2,130
Homeownership share	-0.005 [-0.028, +0.017]	0.641	+0.001	± 10	1,596/2,742
$\ln$ median household income	+0.065** [+0.004, +0.126]	0.038	+0.040	± 10	1,596/2,742
Share aged 65+	-0.004 [-0.014, +0.005]	0.370	-0.002	± 10	1,596/2,742
Prior failed measure $\leq 4y$	-0.018 [-0.104, +0.069]	0.689	-0.016	± 10	1,596/2,743

Notes: Frame: $rd_sample \cap GO$ measures, corpus package v3 (frozen). RBC = robust bias-corrected local-quadratic estimate at h with robust CI (Calonico–Cattaneo–Titiunik, $\rho = 1$); conventional = local-linear, triangular kernel. Bandwidth fixed at $\pm 10pp$ per the reconciliation in Appendix Table A1 (MSE-optimal row and bandwidth-sensitivity curve there). * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$ (robust). Pre-vote covariates; ACS at place grain for municipalities, district grain for schools, county otherwise.

Reading. No pre-vote covariate jumps at the cutoff; the two lagged-issuance indicators and lagged borrowing are smooth, so barely-passed and barely-failed governments were on the same path before the vote.

Table 3: The effect of authorisation on borrowing and building

Outcome	RBC	robust p	Conv.	h	Eff. N (L/R)
Any GO issuance $\leq 6y$	+0.110** [+0.016, +0.204]	0.022	+0.146	± 10	1,596/2,743
Any issuance $\leq 6y$	+0.116** [+0.023, +0.209]	0.014	+0.144	± 10	1,596/2,743
$\ln(1+\text{new-money par p.c.})$	+0.937** [+0.057, +1.818]	0.037	+0.918	± 10	1,537/2,627
$\ln(1+\text{LTD issued p.c.})$, Census survey	+0.878** [+0.146, +1.610]	0.019	+0.827	± 10	1,238/2,128
Same-purpose continuation $\leq 6y$	+0.052 [-0.047, +0.152]	0.302	+0.073	± 10	1,398/2,456
Voter-mode first issuance $\leq 6y$	+0.184*** [+0.089, +0.278]	<0.001	+0.199	± 10	1,596/2,743
No issuance $\leq 6y$	-0.116** [-0.209, -0.023]	0.014	-0.144	± 10	1,596/2,743
Capital outlay p.c. (pay-go bound; schools, post-pre)	+0.464** [+0.090, +0.837]	0.015	+0.377	± 10	996/1,621

Notes: Frame: $\text{rd_sample} \cap \text{GO measures}$, corpus package v3 (frozen). RBC = robust bias-corrected local-quadratic estimate at h with robust CI (Calonico–Cattaneo–Titiunik, $\rho = 1$); conventional = local-linear, triangular kernel. Bandwidth fixed at $\pm 10\text{pp}$ per the reconciliation in Appendix Table A1 (MSE-optimal row and bandwidth-sensitivity curve there). * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$ (robust). Survey row: GFD/IUF, EMMA-independent (survivorship check). Continuation: ballot purpose matched to use-of-proceeds functions; bridge precision 80.0% (blind audit). Pay-go row: full-window cohort (votes to $\sim 2017\text{--}18$, outlay to FY2023).

Reading. A narrow authorisation raises six-year GO issuance by eleven points (robust), roughly doubles per-capita borrowing in disclosure and survey data alike, raises same-purpose project delivery, and moves building nearly one-for-one with borrowing: refusal is not offset by pay-go construction.

Table 4: The fifty-state first stage: rules and observed authorisation

Outcome: voted \$ share (top)	$\beta(\text{strict})$	SE	t	N	Clusters
Pooled (four classes, entity dummies)	+0.0988	0.1165	0.85	13,928	48
School districts	+0.6902	0.1540	4.48	6,857	41
Municipalities	+0.2040	0.0353	5.77	3,176	37
Counties	+0.1078	0.0918	1.17	1,456	43
Special districts	+0.1756	0.0509	3.45	2,439	37
<i>Municipal substitution (general-purpose)</i>					
GO security share (\$)	-0.2767	0.0703	-3.93	4,766	46
Non-chargeable line share (count basis)	-0.0470	0.0209	-2.25	4,600	46

Notes: Entity panel (90,604 local governments), WLS with controls and region FE, state-clustered SEs. RULES PRELIMINARY (pass-1): first-stage associations, not causal estimates. Count-vs-dollar convention for the composition row per Appendix Table A-C1. Townships (proxy rule) excluded.

Reading. The coded rule predicts the observed channel in every adequately powered class, and strict-rule cities substitute away from the voted instrument.

Table 5: The response margin: what refusal buys

Quantity	Value
<i>Panel A: the fate of 100 marginal refusals (2005–19 cohort, $m \leq 5$, $n=422$)</i>	
Re-approved by voters $\leq 4y$	54.3
Returned, not yet converted	13.3
Issued via board or statutory channel	5.2
Issued on pre-existing voter authority	9.0
Extinguished within horizon	18.2
<i>First post-vote authoriser ($m \leq 5$): barely-passed vs barely-failed</i>	
barely-PASSED: voter 48.4%, board 7.9%, none 43.1%	
barely-FAILED: voter 36.7%, board 8.6%, none 54.2%	
<i>Panel B: re-submission (2,680 failed GO measures)</i>	
Hazard, years 1/2/3/4 (%)	26.7 / 22.8 / 15.8 / 12.4
Cumulative return within 4y (%)	58.2
Median time to return (years)	1.02
Returns that pass (%)	61.9
Median amount ratio, return/original	1.000
Purpose categories retained on return (%)	78.2

Notes: Fates are mutually exclusive per 100 barely-refused measures; issuance on pre-existing voter authority draws on authorisations banked before the refused measure and is not evidence of evading it. Panel B: full failed sample; amounts where registries print them ($n=1,354$). Retention from the audited purpose bridge (precision 80.0%).

Reading. Refusal is a pause: most marginal refusals are re-approved at the next election at the same ask, fewer than one in five extinguish, and the board channel is a floor rather than the treatment margin.

Table 6: Where the requirement binds: splits at within-frame medians

Outcome	RBC	robust p	Conv.	h	Eff. N (L/R)
Homeownership: below median	-0.065 [-0.232, +0.103]	0.451	+0.056	± 10	432/863
Homeownership: above median	+0.196** [+0.044, +0.348]	0.011	+0.182	± 10	575/861
difference (above–below)	+0.261**	0.024			
Share 65+: below median	+0.043 [-0.129, +0.215]	0.622	+0.036	± 10	408/781
Share 65+: above median	+0.097 [-0.055, +0.249]	0.210	+0.194	± 10	599/943
difference (above–below)	+0.054	0.646			
Racial-ethnic homogeneity (1–frac.): below median	-0.092 [-0.270, +0.086]	0.311	-0.018	± 10	415/818
Racial-ethnic homogeneity (1–frac.): above median	+0.203*** [+0.059, +0.348]	0.006	+0.232	± 10	592/906
difference (above–below)	+0.295**	0.012			
Median household income: below median	+0.030 [-0.127, +0.187]	0.711	+0.129	± 10	482/802
Median household income: above median	+0.123 [-0.041, +0.288]	0.142	+0.122	± 10	525/922
difference (above–below)	+0.094	0.420			
On-cycle (November, even year)	+0.192** [+0.024, +0.361]	0.025	+0.252	± 10	476/1,005
Off-cycle	+0.072 [-0.039, +0.182]	0.205	+0.099	± 10	1,120/1,738
difference (on–off)	+0.121	0.240			

Notes: Frame: $rd_sample \cap GO$ measures, corpus package v3 (frozen). RBC = robust bias-corrected local-quadratic estimate at h with robust CI (Calonico–Cattaneo–Titiunik, $\rho = 1$); conventional = local-linear, triangular kernel. Bandwidth fixed at $\pm 10pp$ per the reconciliation in Appendix Table A1 (MSE-optimal row and bandwidth-sensitivity curve there). * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$ (robust). Demographic splits run on the 6,255 proper-grain measures (place grain for municipalities, district grain for schools); special districts excluded pending sub-county demographics. Partisanship (precinct-built 577-city panel, all subgroups n.s.) in Appendix Table A-P1.

Reading. The effect concentrates where the consenting public is propertied, older, homogeneous and consulted on-cycle; it is ownership, not income, and institutional, not partisan.

Table 7: Exits and the binding of refusal: the fork against the menu

Outcome	RBC	robust p	Conv.	h	Eff. N (L/R)
Schools, GO issuance	+0.123** [+0.018, +0.228]	0.022	+0.147	± 10	1,289/2,144
Any issuance	+0.132** [+0.027, +0.236]	0.013	+0.150	± 10	1,289/2,144
menu: non-voted \$ share 36.6%; failures re-submitting 0.592					
General-purpose, GO issuance	-0.024 [-0.251, +0.203]	0.837	+0.073	± 10	252/498
Any issuance	-0.028 [-0.249, +0.192]	0.800	+0.045	± 10	252/498
menu: non-voted \$ share 86.1%; failures re-submitting 0.253					
Utilities (special districts), GO issuance	+0.500* [-0.015, +1.016]	0.057	+0.463	± 10	55/101
Any issuance	+0.359 [-0.169, +0.888]	0.182	+0.362	± 10	55/101
menu: non-voted \$ share 73.8%; failures re-submitting 0.426					

Notes: Frame: $rd_sample \cap GO$ measures, corpus package v3 (frozen). RBC = robust bias-corrected local-quadratic estimate at h with robust CI (Calonico–Cattaneo–Titiunik, $\rho = 1$); conventional = local-linear, triangular kernel. Bandwidth fixed at $\pm 10pp$ per the reconciliation in Appendix Table A1 (MSE-optimal row and bandwidth-sensitivity curve there). * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$ (robust). Menu shares from the national corpus (independent of the RD frame). Utilities cells are small; the ordering, not the magnitude, is the cited object.

Reading. Refusal binds exactly where the exit menu is poorest: schools show the discontinuity and the highest re-submission; general-purpose governments, with the richest menu, show none.

Main figures

Figure 1. Issuance at the authorisation threshold

Share with any issuance within six years; 0.5pp bins sized by n; local-linear fits




Figure 1. Issuance at the authorisation threshold

Figure 2. Event study: issuance by year relative to the vote

RBC coefficients with robust 95% CIs; outcome = any new-money issue in relative year k




Figure 2. Event study: issuance by year relative to the vote

Figure 3. The cumulative wedge

Cumulative share with any new-money issue since $k = -2$, $|\text{margin}| \leq 5$; shaded area = the wedge





Figure 3. The cumulative wedge


The geography of consent

Share of local new-money debt (2005–25, \$-weighted) authorised by the voters, all local government classes




By entity class (appendix)


school districts



municipalities



counties



special districts





Figure 4. The consent map: voted share of local new-money dollars

Figure 4/5. Density of the vote margin, by state

0.5pp bins within ± 10 pp of the threshold; the Texas discreteness is visible at the smallest electorates



McCrary log-density: pooled +0.140 (z 2.58); TX +0.206 (z 2.71); all other states n.s. Donut estimates (Table A1) are stable to I

Figure 5. Running-variable density by state

Appendix exhibits (specification battery and landscape detail)

Table A1: Specification battery for the main estimate

Check	Result
Donut $ m > 0.5\text{pp}$	RBC +0.208 [+0.087, +0.329]; conv. +0.195; N 1,502/2,622
Donut $ m > 1.0\text{pp}$	RBC +0.188 [+0.033, +0.343]; conv. +0.186; N 1,409/2,498
Donut $ m > 2.0\text{pp}$	RBC +0.249 [-0.024, +0.521]; conv. +0.199; N 1,199/2,240
Excluding Louisiana	RBC +0.133 [+0.037, +0.229]; conv. +0.162
Cluster by unit (2,480)	conv. +0.146, SE 0.036, z 4.11
Cluster by county (650)	conv. +0.146, SE 0.042, z 3.47
Lee bounds (crosswalk trim 2.52%)	[+0.138, +0.163]
Randomisation inference, $ m \leq 2$	diff. +0.142, $p < 0.0002$ (5,000 permutations)
IK MSE-optimal bandwidth	$h = 1.81\text{pp}$; $\tau = -0.001$ (SE 0.075): slope-dominated on two points of support; the design-based
Alternative horizons 1–6y	Appendix Figure A1a (coefficients rise to the 6y wedge)
Bandwidth sensitivity $h \in [3, 15]$	Appendix Figure A1b (stable from $h \geq 6$)

Notes: Frame: rd_sample $\cap$ GO measures, corpus package v3 (frozen). RBC = robust bias-corrected local-quadratic estimate at h with robust CI (Calonico–Cattaneo–Titiunik, $\rho = 1$); conventional = local-linear, triangular kernel. Bandwidth fixed at $\pm 10\text{pp}$ per the reconciliation in Appendix Table A1 (MSE-optimal row and bandwidth-sensitivity curve there). * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$ (robust). Clustered, Lee, randomisation and IK rows from the committed polish battery (POLISH_RESULTS.md).

Reading. The estimate survives every design-based and inference variant; the one non-rejecting diagnostic (the MSE-optimal local fit) is a power statement, reconciled in the note.

Table A2: Placebo thresholds: no jump where the law draws no line

Outcome	RBC	robust p	Conv.	h	Eff. N (L/R)
50%-regime states, pseudo-cutoff 45 (both sides fail)	+0.143* [-0.004, +0.290]	0.057	+0.012	± 10	786/687
50%-regime states, pseudo-cutoff 60 (both sides pass)	-0.045 [-0.129, +0.039]	0.296	-0.014	± 10	1,874/1,852
CA 66.7% regime, pseudo-cutoff 60 (both sides fail)	-0.068 [-0.394, +0.259]	0.685	-0.007	± 10	69/126
CA 55% regime, pseudo-cutoff 60 (both sides pass)	-0.039 [-0.264, +0.187]	0.736	+0.052	± 10	280/764

Notes: Frame: rd_sample $\cap$ GO measures, corpus package v3 (frozen). RBC = robust bias-corrected local-quadratic estimate at h with robust CI (Calonico–Cattaneo–Titiunik, $\rho = 1$); conventional = local-linear, triangular kernel. Bandwidth fixed at $\pm 10\text{pp}$ per the reconciliation in Appendix Table A1 (MSE-optimal row and bandwidth-sensitivity curve there). * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$ (robust). Pseudo-cutoffs estimated within samples where both sides share the same legal outcome.

Reading. Issuance is smooth through vote shares at which nothing legal changes; the discontinuity appears only at the statutory threshold.

Table A3: The main estimate, state by state

Outcome	RBC	robust p	Conv.	h	Eff. N (L/R)
Texas (50%)	+0.134** [+0.011, +0.256]	0.032	+0.141	± 10	892/1,398
California schools (55%)	+0.059 [-0.154, +0.271]	0.589	+0.105	± 10	261/692
California non-school (66.7%) (small cell)	+0.295** [+0.048, +0.541]	0.019	+0.260	± 10	166/164
Wisconsin (50%)	+0.167* [-0.021, +0.356]	0.082	+0.212	± 10	221/345
Louisiana (50%) (small; parish-fold grain)	-0.396* [-0.796, +0.004]	0.052	-0.276	± 10	40/89
North Carolina (50%) (small cell)	-0.209 [-1.353, +0.934]	0.720	+0.216	± 10	16/53

Notes: Frame: $rd_sample \cap GO$ measures, corpus package v3 (frozen). RBC = robust bias-corrected local-quadratic estimate at h with robust CI (Calonico–Cattaneo–Titiunik, $\rho = 1$); conventional = local-linear, triangular kernel. Bandwidth fixed at $\pm 10pp$ per the reconciliation in Appendix Table A1 (MSE-optimal row and bandwidth-sensitivity curve there). * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$ (robust). Louisiana outcomes fold to the parish; the pooled estimate excluding Louisiana appears in Appendix Table A1.

Reading. Three adequately powered states, three different thresholds, one answer; the small cells are directionally consistent except fold-grain Louisiana.

Table A4: The agenda margin by regime

state (thr)	#districts	proposals/yr	per 100 districts/yr	median amount	on-cycle share
CA (55)	1065	93.5	8.78	\$39,000,000	62.7%
TX (50)	1093	219.2	20.05	\$15,000,000	14.1%
WI (50)	440	48.4	11.00	\$12,600,000	30.5%
MN (50)	392	0	—	—	—

Notes: School GO proposal behaviour 2005–25; MN empty by classification (a data fact). Passage-rate and TX-2019 unbundling panels in the CSV companion (A4b, A4c).

Reading. Under the higher bar districts propose half as often, at 2.6 times the size, four times more on-cycle; pass rates barely move until two-thirds.

Table A5: The menu matrix: non-voted share of new-money dollars

entity type	non-voted \$ share	non-voted issue share	determined \$B
school_district	36.6%	52.2%	509.9
municipal	88.7%	87.6%	578.6
township	62.0%	60.7%	33.7
county	83.1%	86.0%	211.4
special_district	73.8%	51.9%	165.8
authority-class issuer (memo)	97.6%	92.0%	124.7

Notes: National corpus, canonical new-money issues, package v3; determined dollars.

Reading. School districts hold the poorest exit menu; authority-class issuers barely face an electorate at all.

Table A6: The submerged local state

corpus function	project \$B	voted share	lines
public_hospital_facility	197.7	1.89%	2,404
affordable_multifamily_housing	70.0	1.87%	6,017
electric_generation	67.1	0.54%	1,062
natural_gas_utility	50.0	0.05%	232
airport_terminal	48.3	0.25%	540
industrial_development	32.6	0.41%	821
senior_housing	25.8	0.33%	1,087
student_housing_dormitory	12.2	0.07%	383
single_family_homeownership	11.4	0.01%	573
charter_school_facility	6.9	0.00%	596
recycling_waste_facility	6.2	0.34%	429
tax_increment_project_area	5.5	1.38%	852

Notes: Local project functions $\geq \$5B$ voted under 2% of the time; financing mechanics excluded; named examples in text.

Reading. Hospitals, housing, power and airports are financed at scale with essentially no electoral moment.

Table A7: Absolute coalition sizes

entity	cycle	n	votes cast p10/p50/p90	yes-needed p10/p50/p90
school_district	on	829	704 / 4,731 / 38,817	353 / 2,366 / 19,409
school_district	off	4006	270 / 1,293 / 7,133	136 / 647 / 3,567
municipal	on	431	1,101 / 15,217 / 181,750	551 / 7,609 / 90,876
municipal	off	1196	379 / 2,357 / 17,287	190 / 1,179 / 8,644
special_district	on	106	2 / 161 / 2,247	2 / 81 / 1,124
special_district	off	442	2 / 62 / 547	2 / 32 / 274
county	on	173	2,702 / 63,000 / 289,834	1,352 / 31,501 / 144,918
county	off	290	228 / 4,498 / 94,017	115 / 2,250 / 47,009

Notes: TX/WI/LA/NC counts; CDIAC reports percentages only; placeholder rows excluded.

Reading. The same statutory sentence convenes half a million voters in Harris County and two in a developer district.

Table A8: Chargeable share of classified project dollars by authorisation channel

auth mode	\$ch (B)	\$nc (B)	chargeable share
voter	33.8	264.1	11.3%
council_or_board	658.6	439.9	60.0%
statutory	183.8	69.7	72.5%
refunding_no_new_election	73.5	47.4	60.8%
unknown	72.0	34.0	67.9%

Notes: National corpus, classified printed-amount use lines; count-based version in Table A-C1b.

Reading. The voted channel finances the non-chargeable civic core; the unvoted channels carry the chargeable perimeter.

Table A9: Observed voted share by rule regime and entity class

entity	rule	units	voted \$ share	council \$ share	GO security share	non-chargeable share
county	strict	2,375	23.3%	70.9%	49.1%	61.6%
county	non-strict	481	5.4%	90.7%	80.7%	63.1%
municipal	strict	11,037	22.2%	74.1%	37.7%	34.2%
municipal	non-strict	2,305	3.7%	89.2%	72.1%	52.6%
township	strict	5,735	18.0%	80.2%	89.3%	61.0%
township	non-strict	2,483	46.5%	51.8%	97.1%	75.9%
special_district	strict	23,373	33.3%	57.1%	33.0%	30.7%
special_district	non-strict	6,801	13.1%	83.8%	33.4%	18.2%
school_district	strict	12,162	68.4%	30.6%	85.4%	99.4%
school_district	non-strict	616	7.4%	90.6%	45.2%	99.8%

Notes: Dollar-weighted shares from the entity panel (v3); rules PRELIMINARY; townships carry a proxy municipality rule (the reversal is a coding artefact scheduled for pass-2).

Reading. Raw magnitudes behind Table 4: under strict rules two-thirds of school dollars are voted; under lax rules almost none are.

Table A10: The near-miss chain

years since failure k	maj-fail issued $\leq k$ (n at risk)	min-fail issued $\leq k$ (n)	gap
1	1.9% (n=371)	4.6% (n=174)	-2.7%
2	4.9% (n=327)	8.4% (n=155)	-3.5%
3	6.7% (n=326)	10.3% (n=155)	-3.6%
4	7.9% (n=303)	12.8% (n=148)	-4.9%
5	10.6% (n=303)	15.0% (n=147)	-4.4%
6	10.7% (n=270)	16.0% (n=106)	-5.3%
7	13.0% (n=270)	18.9% (n=106)	-5.9%
8	14.5% (n=256)	20.8% (n=96)	-6.4%

Notes: CA failed GO measures; per-cell observability restrictions; new-money issues only.

Reading. The near-miss deficit widens with the window: not a truncation artefact.

Table A11: The rate-cap split (inconclusive)

split	group	conversions	median pass→first
A – within schools (route choice – selection flagged)	schools, 55% return (rate-capped)	118	4.93 (64)
	schools, 66.7% return (uncapped)	9	9.74 (8)
B – cross-class (regime $\not\sim$ entity – CONFOUNDED, flagged)	schools (capped 55)	118	4.93 (64)
	non-schools (uncapped 66.7)	5	12.59 (2)

Notes: 132 conversions; the uncapped cell (14; 9 school) is too thin to test and points against the cap story. Verdict: mechanism neither supported nor excluded.

Reading. Re-approved but unissued stands as a documented fact without an adjudicated mechanism.

Table A12: Demography of blocked majorities (descriptive; California)

covariate	blocked	cleared	sub-50	blocked-cleared (SE)	blocked-sub50 (SE)
renter share	0.370 (n=277)	0.408 (n=1526)	0.351 (n=161)	-0.038 (0.006)	+0.019 (0.010)
Gini (B19083)	0.442 (n=277)	0.450 (n=1526)	0.437 (n=161)	-0.008 (0.002)	+0.005 (0.004)
fractionalization	0.538 (n=277)	0.558 (n=1526)	0.511 (n=161)	-0.020 (0.010)	+0.027 (0.014)
median HH income \$	64,710 (n=277)	68,508 (n=1526)	68,344 (n=161)	-3,798 (1,463)	-3,635 (2,284)
65+ share	0.130 (n=277)	0.130 (n=1526)	0.142 (n=161)	+0.000 (0.003)	-0.012 (0.005)
SAIPE child poverty (schools)	0.155 (n=107)	0.175 (n=696)	0.142 (n=95)	-0.020 (0.010)	+0.013 (0.012)

Notes: Blocked = majority short of the supermajority threshold; within-matched ACS/SAIPE at the D5 grain ladder; DESCRIPTIVE, no causal claim.

Reading. Blocked majorities arise in less affluent, mid-composition school districts: the incidence of the higher bar.

Table A13: Validation of extracted authorisation fields

state	T1: docs w/ date	exact match	±7d	matched→PASSED	T2: voter docs	supported ≤6y	T3: council docs	n
CA	779	62.8%	64.1%	98.0%	849	66.4%	1238	
LA	417	62.8%	63.3%	86.6%	412	86.7%	208	
MA	260	34.6%	40.0%	94.4%	1619	76.9%	756	
MN	209	19.6%	19.6%	85.4%	247	22.7%	1362	
NC	258	65.9%	65.9%	98.8%	0	—	0	
TX	114	73.7%	73.7%	94.0%	7671	76.8%	4454	
WI	900	95.4%	95.6%	87.1%	821	90.5%	443	
pooled	2937	67.9%	68.8%	91.3%	11619	76.2%	8461	

Notes: Election-date matches against independently observed referenda; support and consistency tests as defined in the text. GFD bridge: 2022 public-use issuance matches GFD within 0.5% for 99.9% of bridged units.

Reading. Where registries are complete the extracted fields track the election record at ninety per cent or better; shortfalls sit exactly where registry coverage is known-short.

Figure A1a. Effect by horizon

RBC (solid, with robust CI band) and conventional (dashed)




Figure A1a. RD estimate by issuance horizon

Figure A1b. Bandwidth sensitivity

RBC (solid, with robust CI band) and conventional (dashed)




Figure A1b. Bandwidth sensitivity

Appendix A · Round-4 tables

Table A-C1 · nc_share selection check (classified-line coverage by regime)

Coverage = corpus-active unit (nm_docs>0) with any B3-classified project \$. Weighted coverage weights units by nm_par. Regime = op_referendum_strict (PRELIMINARY labels).

By regime

sample	regime	units	unit coverage	\$ coverage
all classes	strict	12,838	63.5%	82.7%
all classes	non-strict	2,040	77.6%	83.9%
all classes	not codable	3,450	48.3%	80.1%
general-purpose (muni+county)	strict	3,776	68.7%	90.3%
general-purpose (muni+county)	non-strict	1,006	75.1%	84.9%
general-purpose (muni+county)	not codable	1,926	60.9%	92.5%
schools	strict	6,763	66.7%	80.7%
schools	non-strict	263	90.9%	96.1%
schools	not codable	604	3.8%	7.7%

Regime gaps (unit, \$): all classes: 14.1%/1.1%; general-purpose (muni+county): 6.4%/5.4%; schools: 24.1%/15.4% → **UNBALANCED (>5pp) (count-based versions produced below)**.

By state (≥ 20 corpus-active units)

state	units	unit cov	\$ cov
TX	2,131	50.1%	74.1%
IL	1,345	60.7%	75.3%
NY	1,164	65.1%	89.8%
CA	927	46.7%	79.9%
MN	805	79.8%	94.5%
MI	790	72.2%	91.3%
MO	741	70.0%	91.6%
NE	741	72.6%	92.0%
WI	701	85.2%	97.9%
IA	693	74.6%	93.6%
PA	603	3.8%	7.5%
IN	566	4.8%	9.1%
OH	547	77.7%	89.3%
OK	534	56.0%	88.1%
CO	513	45.8%	74.4%
KS	452	88.3%	96.7%
FL	425	42.1%	90.0%
WA	404	75.0%	89.2%
NJ	399	60.7%	80.6%

state	units	unit cov	\$ cov
MA	314	85.0%	96.0%
GA	284	50.7%	71.9%
KY	283	89.8%	98.1%
AR	255	71.0%	84.7%
AL	248	96.0%	97.5%
OR	245	74.3%	89.9%
ND	219	67.1%	93.4%
AZ	202	73.8%	95.0%
CT	164	90.2%	96.9%
SD	159	74.8%	90.1%
TN	157	66.9%	91.3%
MS	155	50.3%	65.2%
SC	129	54.3%	64.7%
UT	128	80.5%	86.1%
LA	114	60.5%	80.0%
NM	113	65.5%	97.3%
NC	112	82.1%	97.0%
MT	111	40.5%	51.2%
ID	91	65.9%	82.6%
VA	75	74.7%	95.6%
ME	60	85.0%	98.0%
WV	40	85.0%	86.9%
MD	37	67.6%	96.0%
NH	36	80.6%	97.5%
RI	34	88.2%	96.8%
NV	30	96.7%	99.8%
WY	29	55.2%	79.2%

Count-based B3 channel sorting (classified LINES, not \$)

auth mode	ch lines	nc lines	chargeable share (count)
voter	15,707	71,961	17.9%
council_or_board	62,488	118,544	34.5%
statutory	7,912	6,727	54.0%
refunding_no_new_election	3,174	4,311	42.4%
unknown	3,840	6,993	35.4%

Count-based N2 (general-purpose): nc line share ~ strict + controls

$\beta(\text{strict}) = \mathbf{-0.0470}$, state-cluster SE 0.0209, t -2.25, n 4,600, clusters 46. (Dollar-based v3 counterpart: -0.087 , t -1.61 .) The text may cite whichever the coverage table justifies; both appear in the appendix.

Table A-P1 · County-grain partisanship (demoted from the text)

County presidential two-party splits of the authorisation effect (GO issuance ≤ 6 y, bw ± 10 ; county Dem share attached to 85.7% of the RD frame): below-median +0.121 (z 2.60) against above-median +0.111 (z 2.30); terciles +0.163 (most Republican) / +0.072 / +0.120 (most Democratic), non-monotone. National first-stage interaction (entity panel, v3): strict $\times$ county Dem -0.65 (t -2.0), descriptive. Caveat: 2020 vintage only, county grain, a coarse proxy for district electorates. The main text's partisanship null cites the precinct-built 577-city panel only, which agrees (all subgroups n.s.).

Appendix H · The honesty record

Moved verbatim from the main text per the round-3 adjudication, updated for P4/P5. The register of informative nulls, pending items and known data limits that the paper carries.

Informative nulls

- Partisanship (three measures, two grains: county presidential, precinct city-footprint, mayoral party): no moderation of the authorisation effect.
- The extensive margin nationally: the coded rule does not predict whether a government borrows at all, only the channel and composition.
- Entity midwifery: no detectable spawning of new issuers in refused units' counties (coarse grain, low power).
- The D2b signal test: no jump in the response at the symbolic 50% line among institutional failures.
- TEL by rule on the big-city subpanel: null, power-limited by design (large cities are the exit-rich class).
- P4 rate-cap split: the banked-authorisation mechanism is neither supported nor excluded (9 uncapped conversions; point direction opposite); the fact stands unexplained.

Pending items, each blocking a specific claim

- Rules human pass-2 (21 cells, with the owner): C2/H2 finals; causal upgrades of the fifty-state first-stage and substitution results; the township rule column (town-meeting states break the municipality proxy).
- Proposition 39 district-level design (state matching-fund controls; CDIAC pre-2005 issuance for the composition margin): journal-version item, not to be started before the conference draft ships.
- B5 audit disagreement review (9 pairs; upgrades the 80% bridge precision).
- Ten-year-plus horizon pass on the re-approved-but-unissued cohort.
- Special-district demographics beyond county proxy (TWDB and shapefile interpolation).
- R1 reform table: secondary__unverified pending the owner's primary-record verification.
- IL and IN vote margins (harvest scoped and paused).

Appendix tables added in round 4

- **Table A-C1 (nc_share selection check):** classified-amount-line coverage

by state and regime, unit- and dollar-weighted (NC_COVERAGE_RESULTS.md). Coverage is regime-UNBALANCED on the unit margin (schools 24.1pp, general-purpose 6.4pp), so every composition exhibit reports a count-based version beside the dollar-based one; the text cites the count-based general-purpose coefficient (-0.047 , $t -2.25$) and both channel-sorting gradients (dollars 11.3/60.0/72.5; counts 17.9/34.5/54.0).

- **Table A-P1 (county partisanship, demoted from the text):** county

presidential two-party splits of the authorisation effect ($+0.121$ against $+0.111$; terciles $+0.163/+0.072/+0.120$, non-monotone) and the national first-stage interaction (-0.65 , $t -2.0$). Caveat: 2020 vintage only, county grain (a coarse proxy for district electorates). The text's partisanship null cites the precinct-built 577-city panel only.

- **Moderation-table note (all D5/H3 tables):** estimates run on the 6,255

proper-grain measures (place for cities, SAIPE/ACS district grain for schools); special districts excluded pending sub-county demographics (TWDB interpolation).

- **Fate-table labels:** "issued anyway" is split into board/statutory

channel (5.2) and pre-existing voter authority (9.0); the latter is not evidence of evasion of the refused measure.

- **Version pin:** the conference draft is frozen on corpus package v3;

package v4 and the B5 rerun are journal-revision items.

- **Pay-go bound cohort:** the outlay row's full-window cohort is votes

through roughly 2017–18 with outlay ending at fiscal 2023.

Known data limits, stated where used

- Louisiana parish-fold grain (outcome mixes measures within a parish).
- Minnesota school-purpose classification gap (empty E4 cell is a

classification fact, not evidence).

- GFD security split (FFC/NG) unreported after 2005 for every entity type and

degenerate for schools before (FFC is 100% of school long-term debt by classification); the corpus security class carries that outcome.

- GFD capital outlay: total only (construction not itemised in the compact;

F-33 not in the repo); no outlay items in the 2024 public-use continuation, so outlay windows end at fiscal 2023.

- Texas BRB counts-unknown placeholder rows ("1-0", 3,188 rows) excluded from coalition-size panels; no RD estimate affected.

- Corpus truncation at 2005 (EMMA era).
- Moderator coverage is within-matched (ACS and SAIPE bridges); matched subsets are not representative of unmatched units.

- Consent-map coverage: ten states carry no determined-mode dollars (an extraction-coverage fact, shown as missing rather than imputed).

Appendix V · Variable inventory by section (§§4–9)

Reference document, 2026-08-24, computed on the current panels (corpus package v3). "Panel" = the 40,924 crosswalked referenda; "RD frame" = $\text{rd_sample} \cap \text{bond_go}$, 11,889 measures. Coverage figures are non-blank counts from the committed files, not estimates. Conditional variables state their own denominator. British spelling; no em dashes.

§4 · Data (identity, treatment definition, validation)

variable	definition	level	source	coverage	notes
unit_id / GID	9-char Census government ID (state2+type1+county3+unit3)	unit	crosswalk	40,924/47,235 referenda (86.6%)	exact tiers ~100% verified; fuzzy tier 95.1% RA-audited
pct_yes	yes share of the vote	referendum	state registries	75% of panel; 100% of RD frame	IL/IN carry results only (no margin)
threshold_centered	margin – applicable threshold	referendum	derived; CA threshold per measure from CDIAC	75% / 100%	the running variable
rd_sample	at a genuine mandatory-ballot statutory cutoff	referendum	rules coding of the 9 states (direct, not the national panel)	23,577 flagged	RD-state cutoffs are bedrock statutes, outside the pass-2 work-list
purpose_class	bond_go / bond_other / tax etc.	referendum	per-state purpose mapping	100%	MN school classification gap flagged
auth_mode_final2	OS-evidenced authorisation mode (voter / council_or_board / statutory)	document	corpus extraction	93.7% of docs determined (v3)	validated: date match 67.9% pooled, 95.4% WI; matched→passed 91.3%
op_referendum_strict (+ ordinal, threshold_num)	strict rule, state × entity × purpose × year	state × entity × year	panel	89% of panel rows; muni cell codable in 37 states	PRELIMINARY pass-1; 21-cell pass-2 with owner

§5 · Landscape and fifty-state first stage

Grain: one row per local government (entity panel, 90,604 units) or per document (corpus).

variable	definition	level	source	coverage	notes
nm_docs, nm_par	new-money issues and par, 2005–25, canonical per issue	unit	corpus v3	78,672 typed nm issues; 46 states ≥50 docs	5 small states below gate (DC/DE/HI/VT/WY); 5 legacy states 72–87% flag coverage (MN/MA/MO/MD/ID)
sec_go_sh ... sec_sptax_sh	par shares by security pledge class	unit	corpus	where nm_par>0	GO / revenue / lease / special_tax

variable	definition	level	source	coverage	notes
voted_sh_par, council_sh_par, statutory_sh_par	mode shares of determined nm \$	unit	corpus	13,875 units with a determined split	the first-stage outcome
nc_share_project	non-chargeable share of classified project \$	unit	B3 doc flags (118-label map, 0 unmapped)	~18% of units	classified printed-amount lines only
ballot purposes (M2 panel A)	normalised keyword categories	referendum	9 registries	19,600 bond measures	stand-in for a formal B1 taxonomy (flagged)
votes cast / yes-needed (M3)	counts and threshold $\times$ total	referendum	TX/WI/LA/NC registries	7,473 GO measures	CA has % only; 3,188 TX placeholder rows excluded; LA also carries voters_qualified/voted
GFD characteristics	pop or enrolment, revenue, own-source, property tax, LTD out	unit, latest yr ≥ 2012	national GFD	100% of entity panel	specials lack Population (revenue-size proxy in N-suite)
ACS covariates	homeownership, 65+, fractionalisation, median income	place (munis) / county (rest)	ACS5-2019 national pulls	~100%	county layer uniform; place upgrade for cities
county_dem2p_2020	county Dem two-party share	county	MEDSL	99–100%	2020 only
FOG institutions	form of gov, initiative, referendum, partisan, districts	city	FOG panel	~570 big cities	subpanel only
TEL stringency; ACS-2022; city Dem; mayor party rule (per class)	big-city extras strict for the unit's entity class	city state $\times$ entity	municipal-analysis panels rules panel	~570 cities each county 94% · schools 93% · specials 79% · munis 68% · townships 50% (proxy)	2013-vintage TEL, no time variation township column = pass-2 Part C

§6 · The RD core (frame: 11,889 GO measures)

variable	definition	level	source	coverage in RD frame	window
issued_6y	any issuance $\leq 6y$	referendum	corpus link	100%	(0, +6y]
GO issuance $\leq 6y$	issued_6y $\wedge$ go_share >0	referendum	corpus	100% (go_share observed for 64%, else 0-issue)	headline outcome
nm_par_6y, ln_par_pc_6y	new-money par; ln(1+par p.c.)	referendum	corpus + GFD denominator	100% / 86%	denominator = gfd_pop, enrolment for schools (65% of frame are schools)
ln_gfd_ltd_pc_6y	survey LTD issued p.c. (EMMA-independent row)	referendum	GFD + IUF FY23/24	73%	fiscal [y+1,y+6]; IUF validated 99.9%

variable	definition	level	source	coverage in RD frame	window
capital outlay p.c. (pay-go bound row)	GFD To- tal_Capital_Outlays p.c.	referendum	GFD	83% with a 6y window (8,489/10,244 with denominator)	ends FY2023 (IUF has no outlay); total outlay only
ev_m2...ev_p5	any nm issue in relative year k	referendum	corpus	100%	event study
balance covariates	pre-vote GFD pop, revenue p.c., debt p.c., taxes, enrolment, prior failure	referendum	GFD	94–98%	0/7 imbalanced
McCrary/density inputs	binned margin counts	bin	derived	full frame	TX-driven excess

§7 · The response margin (denominators are conditional)

variable	definition	level	source	coverage	notes
resubmitted_4y hazard, KM	same-unit+class return $\leq 4y$	failed measure	panel sequences	2,680 failed GO measures (the denominator)	risk sets censored at registry end
return pass / amount ratio	outcome of the return	return	registries	1,418 returns; 1,354 with amounts	amounts CA/TX/WI
fate table categories	converted / returned / issued-anyway / extinguished	failed measure	corpus + sequences	1,233 (2005-19 full windows); 422 close	mutually exclusive
first-issue authoriser	mode of first post-vote nm issue	referendum	corpus	close window n=2,234	transition matrix
continuation (B5)	window doc sharing the ballot's purpose category	measure × do	EN2CAT bridge	3,854 close categorisable measures	bridge precision 80.0% / recall 88.9% (blind audit); re-runs on v4
bundle re-composition P3 chain	categories kept/dropped on return pass → first-issue timing; 8y issuance	failure → return pair conversion	bridge corpus	1,129 pairs 132 CA conversions (118 capped / 14 uncapped)	uncapped cell too thin (P4 not confirmed)

§8 · Heterogeneity (moderators; split at within-frame medians)

variable	definition	level (grain)	source	coverage in RD frame	notes
acs_homeown/share65+ grain/moderators	place moderators	place SD county	2,103 4,152 5,329	ACS5 2010/2019 (pre-vote vintage rule)	97% attached proper-grain-only splits use the 6,255 place+SD rows
sd_childpov_rate	district child poverty (income proxy)	school district	SAIPE name-bridge	35% of frame (4,153 schools)	match rates vary (MN 91% ... CA 44%); within-matched validity superseded by proper grain where available
county mod-erators	65+, frac, income (first-pass proxy)	county	CC-EST SAIPE	+ 97.4%	
county_dem2p (pre-vote)	partisanship	county	MEDSL	85.7%	null result

variable	definition	level (grain)	source	coverage in RD frame	notes
city Dem share, mayor party	large-city par-tisanship	city (precinct-built)	municipal-analysis	1,932 municipal mea-sures in 577-city panel; 178 within ± 10	replicate-the-null only
Gini (D6)	B19083	place/SD/county	new CA pull (2010+2019) dates	CA cells (blocked-majority analysis) 100%	D6 is CA-only by construction
election timing	on-cycle = Nov of even year	referendum			
entity class (fork)	census_type	unit	crosswalk	100%	schools 7,690 · GP 2,623 · specials 1,576

§9 · Agenda margin and politics

variable	definition	level	source	coverage	notes
proposals /100 districts/yr; on-cycle share	dis-ask; E4 regime comparison	state×year	registries + GFD district counts	CA/TX/WI (MN = classification gap)	
pass rates by regime	E3	regime	registries	11,888 with margins 2014–25	invariance at 50 vs 55; breaks at 66.7 reporting-mandate caveat
props/election; multi-prop share	TX-2019 un-bundling	district×election	ENBRB		the DiD-at-fixed-support
W1/W2 wedge cells	issuance by support band × regime	band×state×panel group		50–55 band: CA 187 / TX+WI 653	
reform events	attempts/outcomes since 1990	state×event	web compilation	11 rows	secondary_unverified (owner)
Prop 39 / M56 DiD outcomes	ln per-pupil (per-capita) LTD issued	state×year	national GFD 1997–2013	schools ≥ 20 districts/cell, 39–48 states	state-mean grain; district-level upgrade deferred

The three load-bearing coverage facts a referee will probe

1. **RD outcomes are 100%/86% covered in-frame** (binary/intensive) and the survey replication row (73%) is symmetric across the cutoff; crosswalk selection is Lee-bounded. 2. **Moderator grains are honest**: place for cities, district for schools (SAIPE/ACS-SD), county fallback (all specials); splits are within-matched and labelled; nothing is imputed. 3. **National corpus coverage post-v3**: 46 states above the 50-doc gate; the residuals are five thin small states and five legacy partial-fill states, both listed; rule coding 68–94% by class with the township proxy flagged.