Who Must Agree

Bond authorisation rules as coalition requirements

Empirical sections and exhibits — working conference draft

Compiled 25 August 2026 from the committed analysis record (corpus package v3, frozen).

*Working conference draft. Every number traces to a committed results file (script-to-results map in `analysis/ANALYSIS_REVIEW.md`). **[PENDING]** marks claims held on a named data pass. The draft is frozen on corpus package v3. Sections 1 to 3 (theory and institutional framework) are maintained separately.*

Reading paragraphs under exhibits are drafting aids, marked strippable for submission.

4 · Data: observing the consent requirement end to end

The question of this paper is what a voter-approval requirement for local borrowing actually does. Answering it requires seeing three things at once, and no single dataset contains them. The first is the **rule**: which governments must ask voters before borrowing, and what share of the vote they need. The second is the **asking**: what governments put on the ballot, and whether it passed. The third is the **borrowing**: what was ultimately issued, for what purpose, and, critically, on whose authority. This section describes how each is measured. Readers familiar with municipal finance data can skim to Table 1, which summarises the samples; a full variable inventory with sources and coverage is Appendix V.

The referendum panel: 47,235 ballot measures. Nine states publish administrative registries of local bond and tax elections: California (the CDIAC database), Texas (Bond Review Board), Wisconsin (Department of Public Instruction), Louisiana (Secretary of State), North Carolina (State Board of Elections), Massachusetts (Division of Local Services), Minnesota, Illinois and Indiana. Compiled, these give 47,235 measures. Each measure is then matched ("crosswalked") to its government in the Census of Governments, so that election outcomes can be joined to that government's later borrowing: 40,924 measures (86.6%) resolve to a Census identifier, with the exact-match tiers verified at close to 100% and the fuzzy-match tier independently audited at 95.1% accuracy.

Not every measure is useful for causal analysis. The design of Section 6 needs measures where a statute, not a local choice, sets a pass threshold and makes the ballot mandatory. Restricting to such measures leaves 23,577. Restricting further to **general-obligation (GO) bond measures**, bonds backed by the government's taxing power, which is the instrument statutes typically make voters approve, leaves the analysis frame: **11,889 GO measures** (Texas 8,062, California 2,189, Wisconsin 999, Louisiana 361, North Carolina 278). These sit at three statutory thresholds: 50% (Texas, Wisconsin, Louisiana, North Carolina), 55% (California school measures) and 66.7% (California non-school measures), with each Californian measure's threshold taken from CDIAC. Two derived variables matter throughout: `pct_yes`, the yes share of the vote, and the **running variable**, `pct_yes` minus the threshold that applied to that measure, so that zero always marks the pass/fail line.

The outcome of record: the official-statement corpus. When a local government sells bonds publicly, it files an official statement, a disclosure document describing the issue. The corpus assembled here contains **258,762 official statements** covering 2005 to 2025, all fifty states and 43,030 issuers. From each document are extracted the par amount, the security class (GO, revenue, lease, special tax), the use-of-proceeds lines classified into 118 functional activities, and the variable the paper turns on: the **authorisation mode the document itself evidences**. An official statement states the legal authority for the issue, so each issue can be coded as authorised by voters (usually with the election date cited), by the governing board, or by statute. This is the richest available measure of what was borrowed, for what, and on whose authority, and every headline claim about provision (issuance, channel, purpose, project continuation) rests on it.

Census survey data enter in exactly two supporting roles, each a single table row. The Government Finance Database debt items (1967 to 2023, 2.1 million local unit-years) provide a survivorship check: the corpus only sees publicly sold debt, the survey also counts bank loans and private placements, so finding the same effect in both rules out the concern that refused governments simply borrow off-market. The survey capital-outlay items provide a bound on the one channel the corpus cannot see by construction: building without borrowing (the pay-go bound, Section

Table 1: Sample and summary statistics

State (threshold)	N	Mean yes	Passed	Median amount	Median electorate
CA (55%)	1,715	62.423	0.809	\$44.0m	–
CA (66.7%)	466	66.442	0.509	\$20.0m	–
LA (50%)	361	59.796	0.770	–	1,357
NC (50%)	278	65.843	0.914	–	30,568
TX (50%)	8,062	68.849	0.792	\$15.0m	1,397
WI (50%)	999	54.070	0.661	\$12.4m	2,988
All compiled referenda	47,235				
Matched to a Census government	40,924				
At a genuine mandatory-ballot requirement	23,577				
GO regression-discontinuity frame	11,889				

Panel B: the official-statement corpus (package v3)

Documents: 258,762
Issuers: 43,030
Determination rate: 0.937
Mode shares (docs): voter / board / statutory: 0.195 / 0.666 / 0.063
New-money par: schools / municipal / county / special / township (\$B): 524 / 606 / 220 / 139 / 35

Notes: Panel A: the referendum frame by state and threshold (cells $\geq$30 measures); amounts where registries print them (CA/TX/WI); electorate = votes cast where counts exist (TX/WI/LA/NC; counts-unknown placeholder rows excluded). Restriction cascade beneath. Panel B: corpus package v3.

Reading. The frame spans three statutory thresholds; the corpus supplies the authorisation mode for 93.7 per cent of a quarter-million documents.

6). Neither use displaces the corpus as the outcome of record.

Validating the extracted authorisation mode. Because the authorisation mode is read out of documents rather than taken from an administrative source, it is checked against the election record before use. Where an official statement cites an election date, that date matches an independently observed referendum for the same government 67.9% of the time pooled, and 95.4% in Wisconsin, the one state whose registry is known to be complete; conditional on a match, 91.3% of cited elections are passed measures (98.0% in California, 98.8% in North Carolina). The shortfalls sit exactly where registry coverage is known to be short: Minnesota's registry begins in 2020, and Massachusetts records votes at month grain. A related check runs the other way: 76.2% of voter-mode new-money documents have a passed referendum within the previous six years in the records, and the gap is driven by authorisations older than the observation window (Texas districts routinely issue against decade-old voter authorisations in series). On the survey side, the 2022 public-use issuance item matches the Government Finance Database within 0.5% for 99.9% of bridged units. The full validation table is Appendix Table A11.

The rules panel. The fifty-state analysis of Section 5 needs a coding of each state's rule: for every state, entity type (county, municipality, township, school district, special district), borrowing purpose and year, does the default path to GO debt run through a mandatory ballot, and at what threshold? The working definition of a **strict rule** is that a pre-issuance ballot referendum at the polls is the mandatory default path. Rules that demand a town-meeting vote, a referendum only if petitioned, or a vote only above a debt limit or for certain purposes are coded non-strict, because in each case a government can normally borrow without convening an electorate at the polls. **[PENDING]** The current panel is a first-pass machine coding, cross-validated at 78% (29 of 37 codable states) against an independent hand coding; a 21-cell human verification pass is in progress, with archived statutory texts. Until it lands, every rule coefficient in this paper is presented as a first stage, the association between the coded rule and observed authorisation behaviour, never as a causal estimate. Each table note repeats the convention.

5 · The national landscape: rules, menus and the fifty-state first stage

This section widens the lens from nine states to all fifty. It proceeds in four steps. Section 5.1 builds the national panel and defines its outcome variables. Section 5.2 documents three facts about the institutional landscape that motivate everything after. Section 5.3 asks the first-stage question: does the coded rule show up in observed authorisation behaviour (Table 5)? Section 5.4 asks what governments do instead where the rule binds (Table 6), and Section 5.5 states honestly what reform events can and cannot yet identify.

5.1 A panel of every local government

The unit of analysis is the individual local government. The panel covers **90,604 units**: every county, municipality, township, school district and special district in the national Government Finance Database universe. Each unit carries four blocks of variables. First, the coded rule for its state and class (Section 4). Second, its borrowing 2005 to 2025 from the corpus: the count and par of new-money issues, and for each issue its security class and authorisation mode. The headline outcome, the **voted dollar share**, is the share of a unit's determined new-money dollars that the documents evidence as voter-authorised. Third, its borrowing 2005 to 2023 from the survey. Fourth, demographic and fiscal controls: population (for special districts, which lack a population concept, ln total revenue serves as the size control), homeownership, the share aged 65 and over, racial fractionalisation, median household income, and county presidential partisanship; place-grain values for cities, county-grain otherwise. A subpanel of roughly 570 large cities adds institutional detail: form of government, tax-and-expenditure-limit stringency, city-level partisanship and mayoral party.

Two composition outcomes recur and deserve definition. The **GO security share** is the fraction of a unit's new-money dollars issued as general-obligation debt rather than revenue bonds, leases or other instruments that typically need no vote. The **non-chargeable share** is the fraction of classified project dollars flowing to functions that cannot be billed to users (schools, roads, parks, jails), as opposed to chargeable ones (water systems, airports, hospitals) whose revenue streams support unvoted revenue bonds. Because line-level dollar coverage differs across rule regimes, the non-chargeable share is measured on two bases, dollars and classified line counts; Appendix Table A4 documents the imbalance and the count basis is the one the text cites where they differ.

5.2 Three facts about the landscape

Fact 1: different kinds of government hold radically different menus of exits from the voted channel. Table 2 tabulates, for the national new-money corpus (78,672 canonical issues), the share of determined dollars authorised without a vote, by entity type. School districts route only **36.6%** of new-money dollars through non-voted channels. At the other end, municipalities route **88.7%**, counties 83.1%, and authority-class issuers (housing, hospital and utility conduits)

97.6%; townships (62.0%) and special districts (73.8%) sit between. The reason is visible in the full matrix of security class by authorisation mode: a school district's menu is essentially GO-or-nothing, while a city can finance most of what it does through revenue bonds, leases and conduit authorities that never face an electorate. "Who must agree" is not one institution. It is a menu that varies by the kind of government doing the asking.

Table 2: The menu matrix: non-voted share of new-money dollars

entity type	non-voted $ share	non-voted issue share	determined $B
school_district	36.6%	52.2%	509.9
municipal	88.7%	87.6%	578.6
township	62.0%	60.7%	33.7
county	83.1%	86.0%	211.4
special_district	73.8%	51.9%	165.8
authority-class issuer (memo)	97.6%	92.0%	124.7

Notes: National corpus, canonical new-money issues, package v3; determined dollars.

Reading. School districts hold the poorest exit menu; authority-class issuers barely face an electorate at all.

Fact 2: voters are shown the civic core; the chargeable perimeter never votes. Classifying ballot purposes across 19,600 bond measures, K-12 schools alone account for 38.2% of everything put before voters, followed by water and sewer, roads, parks and fire protection (Table 3, Panel A). Panel B lists the corpus functions whose dollars are voted on less than 2% of the time: public hospitals ($197.7B of local project dollars), multifamily housing ($70.0B), electric generation ($67.1B), gas utilities, airport terminals. The contrast is concrete. A $0.84B Harris County Hospital District line, San Francisco's multifamily housing revenue programme and Energy Northwest's generation debt faced no ballot; Los Angeles Unified's $9.0B school ask in 2024 did. At the channel level the sorting is stark: the voted channel carries **11.3%** chargeable dollars, the board channel 60.0% and the statutory channel 72.5% (in classified line counts: 17.9%, 34.5% and 54.0%, the basis robust to coverage; Appendix Table A5 gives the dollar matrix by channel). What the consent requirement governs, overwhelmingly, is the class of goods that cannot be charged to users, precisely the class the theory says it should.

Table 3: The submerged local state

corpus function	project $B	voted share	lines
public_hospital_facility	197.7	1.89%	2,404
affordable_multifamily_housing	70.0	1.87%	6,017
electric_generation	67.1	0.54%	1,062
natural_gas_utility	50.0	0.05%	232
airport_terminal	48.3	0.25%	540
industrial_development	32.6	0.41%	821
senior_housing	25.8	0.33%	1,087
student_housing_dormitory	12.2	0.07%	383
single_family_homeownership	11.4	0.01%	573
charter_school_facility	6.9	0.00%	596
recycling_waste_facility	6.2	0.34%	429
tax_increment_project_area	5.5	1.38%	852

Notes: Local project functions ≥$5B voted under 2% of the time; financing mechanics excluded; named examples in text.

Reading. Hospitals, housing, power and airports are financed at scale with essentially no electoral moment.

**Fact 3: the same legal sentence assembles coalitions that differ by five orders of

magnitude. Where vote counts exist (Texas, Wisconsin, Louisiana and North Carolina; 7,473 GO measures, since California reports percentages only), the median school bond was decided by **768** yes-votes, the median off-cycle city measure by 1,179, and the median special-district election by **34**. Under the identical Texas 50% rule, Harris County's road programme required 511,375 yes-votes of 1,022,748 cast, while a developer municipal utility district was authorised by **two votes of two cast** (Table 4). The coalition a referendum requirement demands is not a constant of the law but a variable of the electorate the law convenes. Section 8 returns to this observation. (Texas rows carrying the registry's counts-unknown placeholder, 3,188 rows recorded "1-0", are excluded and documented; no RD estimate is affected.)

Table 4: Absolute coalition sizes

entity	cycle	n	votes cast p10/p50/p90	yes-needed p10/p50/p90
school_district	on	829	704 / 4,731 / 38,817	353 / 2,366 / 19,409
school_district	off	4006	270 / 1,293 / 7,133	136 / 647 / 3,567
municipal	on	431	1,101 / 15,217 / 181,750	551 / 7,609 / 90,876
municipal	off	1196	379 / 2,357 / 17,287	190 / 1,179 / 8,644
special_district	on	106	2 / 161 / 2,247	2 / 81 / 1,124
special_district	off	442	2 / 62 / 547	2 / 32 / 274
county	on	173	2,702 / 63,000 / 289,834	1,352 / 31,501 / 144,918
county	off	290	228 / 4,498 / 94,017	115 / 2,250 / 47,009

Notes: TX/WI/LA/NC counts; CDIAC reports percentages only; placeholder rows excluded.

Reading. The same statutory sentence convenes half a million voters in Harris County and two in a developer district.

5.3 The first stage: does the rule show up in behaviour?

The first-stage question is simple: comparing governments of the same class under strict and non-strict rules, is more of their borrowing voter- authorised where the statute says it must be? Table 5 reports the regression: the voted dollar share on the strict-rule indicator, estimated by weighted least squares with region fixed effects, the controls of Section 5.1 and standard errors clustered by state, pooled and then class by class.

Table 5: The fifty-state first stage: strict rules and the voted share of borrowing

	(1) Pooled	(2) Schools	(3) Municipalities	(4) Counties	(5) Special districts
Strict referendum rule	$+0.099$	$+0.690^{***}$	$+0.204^{***}$	$+0.108$	$+0.176^{***}$
	(0.117)	(0.154)	(0.035)	(0.092)	(0.051)
Controls	Yes	Yes	Yes	Yes	Yes
Region fixed effects	Yes	Yes	Yes	Yes	Yes
Entity-class dummies	Yes	–	–	–	–
Observations	13,928	6,857	3,176	1,456	2,439
State clusters	48	41	37	43	37

Notes: Dependent variable: the voted share of the government's determined new-money dollars, 2005–25, measured from the official-statement corpus (package v3). Columns (2)–(5) estimate the class-specific coefficient on its own subsample. Townships are excluded (they carry a proxy municipal rule; adding them leaves the pooled coefficient at +0.090, t 0.74). Entity panel (90,604 local governments; estimation samples below are the corpus-active units for each outcome). Weighted least squares; Census-region fixed effects; controls: ln population (ln total revenue for special districts), homeownership, share 65+, racial fractionalisation, ln median household income, county Democratic presidential share 2020. Standard errors clustered by state in parentheses. The strict-rule indicator is the PRELIMINARY pass-1 coding (verification pass in progress), so coefficients read as first-stage associations, not causal estimates. * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.

The rule shows up, and it shows up most where the theory says it must. Under a strict rule the voted share of new-money dollars is higher by **+0.690 for school districts (t = 4.5)**, +0.204 for municipalities (t = 5.8) and +0.176 for special districts (t = 3.5); the county coefficient (+0.108) is not significant. The pooled coefficient is small and insignificant, and column (1) is reported to show why class-by-class estimation is the right object: the pooled number averages a large school effect with small general-purpose effects across very different subsample compositions. Raw shares tell the same story without any regression: 68.4% of school dollars are voted under strict rules against 7.4% under non-strict ones; 22.2% against 3.7% for municipalities; 23.3% against 5.4% for counties (Appendix Table A6 reports the full grid).

Figure 1 maps the resulting geography of consent for the 46 states passing the coverage gate: the observed voted share of local new-money debt ranges from roughly 80% in Oklahoma and two-thirds in Texas to under 3% in New York, Pennsylvania, Tennessee and Kentucky, the states whose statutes require no local ballot.

One reversal in the class table is itself diagnostic. Townships show 18.2% voted under nominally strict rules against 46.6% under non-strict ones, the opposite gradient to every other class. Townships carry a proxy rule in the panel (the municipal coding of their state), and New England towns borrow by town-meeting vote in states coded non-strict for cities. The panel is signalling that the coding needs a township class; the anomaly is flagged rather than absorbed, and the township column is on the verification worklist. [**PENDING: rules pass-2.**]

5.4 Substitution: what governments do instead

If a strict rule makes the GO instrument costly, governments with alternatives should shift towards them. Table 6 tests this on the composition of borrowing, and then asks whether the rule changes the amount of borrowing at all.

Table 6: Substitution away from the voted instrument, and the quiet extensive margin

	(1) GO Pooled	(2) GO GP	(3) NC $ GP	(4) NC lines GP	(5) Any issue All	(6) ln issued All
Strict referendum rule	-0.155*	-0.277***	-0.087	-0.047**	+0.015	-0.232
	(0.086)	(0.070)	(0.054)	(0.021)	(0.029)	(0.217)
Controls	Yes	Yes	Yes	Yes	Yes	Yes
Region fixed effects	Yes	Yes	Yes	Yes	Yes	Yes
Observations	14,273	4,766	3,340	4,600	56,541	56,541
State clusters	48	46	46	46	48	48

Notes: Dependent variables: (1)–(2) GO = share of new-money dollars issued as general-obligation (GO) security; (3) share of classified project dollars in non-chargeable functions (services that cannot be billed to users); (4) the same composition measured in classified use-of-proceeds line counts, the basis robust to the coverage imbalance documented in Appendix Table A-C1 and therefore the estimate the text cites; (5) Any issue = linear probability of any corpus new-money issue 2005–25; (6) ln issued = log survey long-term debt issued per capita 2005–23 (Government Finance Database). GP = general-purpose governments (municipalities and counties); schools are near-degenerate for the non-chargeable outcome and dilute the pooled composition rows. Entity panel (90,604 local governments; estimation samples below are the corpus-active units for each outcome). Weighted least squares; Census-region fixed effects; controls: ln population (ln total revenue for special districts), homeownership, share 65+, racial fractionalisation, ln median household income, county Democratic presidential share 2020. Standard errors clustered by state in parentheses. The strict-rule indicator is the PRELIMINARY pass-1 coding (verification pass in progress), so coefficients read as first-stage associations, not causal estimates. * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.

Cities under strict rules substitute away from the voted instrument. At the unit grain with

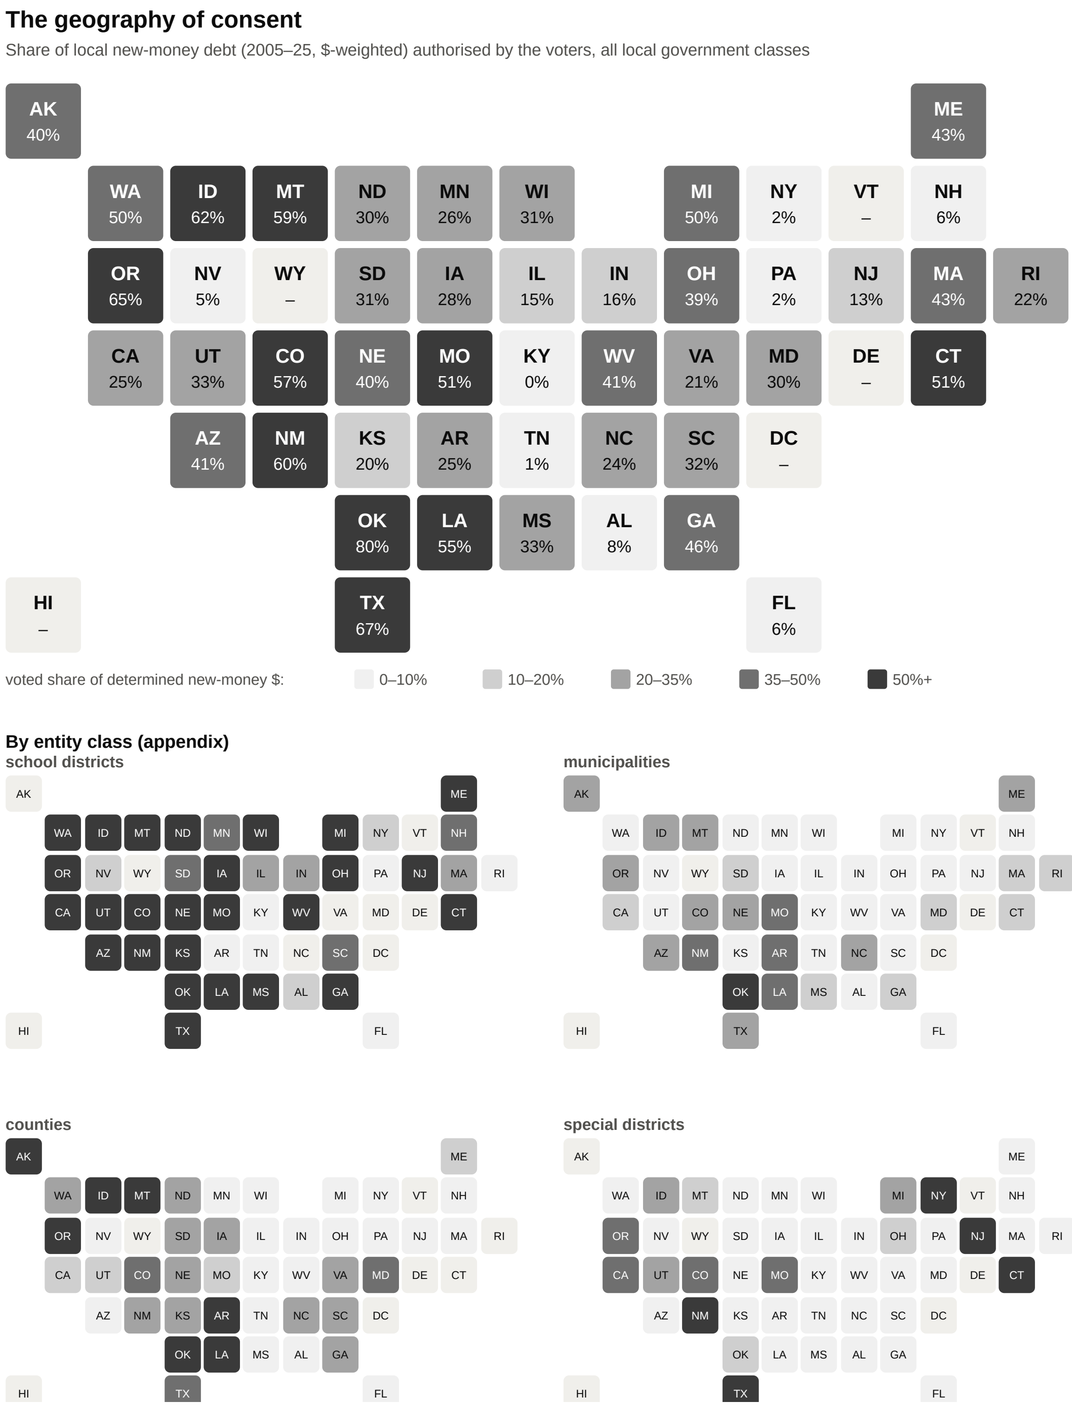

Figure 1: The geography of consent. Voted share of local new-money dollars by state, official-statement corpus 2005–25, package v3. States below the 50-document coverage gate (DC, DE, HI, VT, WY) are shown as no-data.

full controls, general-purpose governments in strict states carry a GO security share **27.7 points lower** (t = −3.9) than their counterparts in non-strict states (column 2). Their project composition shifts in the same direction: the non-chargeable gap is −8.7 points (t = −1.6) measured in dollars and **−4.7 points (t = −2.3)** measured in classified line counts (columns 3 and 4), the count basis being the one robust to the coverage imbalance documented in Appendix Table A4. A state-level fixed-effects version of the same test gives −0.162 (t = −1.83). This is the fifty-state generalisation of Fact 2: where the GO instrument requires assembling a coalition, governments that can finance through chargeable, unvoted instruments do so.

The extensive margin, by contrast, is quiet. Whether a government borrows at all over 2005 to 2025 (column 5) and how much it issues in survey totals (column 6) are not significantly associated with the rule. The rule moves the channel and the composition of borrowing, not its existence, exactly the pattern the regression-discontinuity results of Section 6 rationalise: the requirement prices delay and channel, not denial.

Institutional interactions on the big-city subpanel (tax-and-expenditure limits) and national demographic interactions are reported in Appendix Table A7; none is significant except a negative interaction with county Democratic share (t = −2.0), which is county-grain, 2020-vintage and read as descriptive. [**PENDING: all rule coefficients in this section are first-stage or descriptive until the rules verification pass.**]

5.5 What reform events can and cannot show

The strongest available national causal design is California's Proposition 39 (November 2000), which cut the school threshold from two-thirds to 55% while leaving every other Californian government and every other state untouched. It does not yet deliver a verdict at the state-mean grain. The schools-only difference-in-differences is +0.45 log points (permutation p = 0.15 across 49 placebo states) with a visible pre-trend, and the sector-differenced triple-difference is approximately zero. Two data facts discipline the design rather than the theory: the survey's pre-2005 security split classifies all school long-term debt as full-faith-credit, so the composition shift Proposition 39 should produce is invisible in Census data; and California's state matching-bond waves (Propositions 47 and 55, 2002 and 2004) confound the totals margin. [**PENDING: the publishable version is a district-level two-way fixed-effects design with enrolment weights, Conley-Taber inference, matching-fund controls and CDIAC issuance data; a journal-version item.**] The reform record accordingly appears in Section 9 as evidence on the politics of the rule, and the causal weight of the paper rests on the within-state designs of Section 6.

6 · Authorisation at the margin: regression-discontinuity evidence

The design, in plain terms. A government that wins its referendum with 50.4% and one that loses with 49.6% are, in every respect that matters, the same government on the same day; the only thing that differs is which side of the statutory line the vote landed. Comparing many such near-winners with near-losers therefore isolates the causal effect of authorisation itself. Formally, for the 11,889 GO measures at statutory thresholds, the outcome is regressed on the running variable (the vote share minus the threshold) separately on each side of zero, local-linear with a triangular kernel within a ±10 percentage-point bandwidth. Inference follows the robust bias-corrected (RBC) convention of Calonico, Cattaneo and Titiunik: each table reports the RBC estimate with its robust confidence interval and p-value, the conventional local-linear estimate beside it, and the effective sample on each side of the cutoff.

Balance. The design's premise is that near-winners and near-losers are comparable before the vote. Table 7 checks it: none of seven pre-vote covariates (lagged issuance, population, enrolment, homeownership, income, age structure, prior measure count) jumps at the cutoff (maximum $|t| = 1.36$).

Table 7: Covariate continuity at the authorisation threshold

Outcome	RBC	robust p	Conv.	h	Eff. N (L/R)
Any new-money issue, year $t-2$	+0.030 [-0.028, +0.089]	0.314	+0.007	±10	1,596/2,743
Any new-money issue, year $t-1$	+0.021 [-0.036, +0.079]	0.468	+0.022	±10	1,596/2,743
ln(1+LTD issued p.c., $t-3..t-1$)	+0.530 [-0.256, +1.315]	0.186	+0.412	±10	1,495/2,553
ln population	+0.320 [-0.521, +1.160]	0.456	+0.142	±10	252/497
ln enrolment (schools)	-0.035 [-0.343, +0.272]	0.822	-0.001	±10	1,285/2,130
Homeownership share	-0.005 [-0.028, +0.017]	0.641	+0.001	±10	1,596/2,742
ln median household income	+0.065** [+0.004, +0.126]	0.038	+0.040	±10	1,596/2,742
Share aged 65+	-0.004 [-0.014, +0.005]	0.370	-0.002	±10	1,596/2,742
Prior failed measure ≤4y	-0.018 [-0.104, +0.069]	0.689	-0.016	±10	1,596/2,743

Notes: Frame: rd_sample ∩ GO measures, corpus package v3 (frozen). RBC = robust bias-corrected local-quadratic estimate at h with robust CI (Calonico–Cattaneo–Titiunik, $\rho = 1$); conventional = local-linear, triangular kernel. Bandwidth fixed at ±10pp per the reconciliation in Appendix Table A1 (MSE-optimal row and bandwidth-sensitivity curve there). * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$ (robust). Pre-vote covariates; ACS at place grain for municipalities, district grain for schools, county otherwise.

Reading. No pre-vote covariate jumps at the cutoff; the two lagged-issuance indicators and lagged borrowing are smooth, so barely-passed and barely-failed governments were on the same path before the vote.

Density. One threat is specific to this setting: governments time proposals to win, so there is excess mass just above the cutoff (McCrary $\theta = +0.24$, z = 4.6). Figure 2 shows where it lives: it is Texas-specific (TX +0.21, z 2.7; every other state approximately zero), carries the signature of endogenous proposal timing documented since Cellini, Ferreira and Rothstein (2010), and is itself evidence on the agenda margin taken up in Section 9. Identification survives it: donut estimates that drop the contested region around the cutoff are stable to larger (+0.19 to +0.20, z 3.7 to 5.3; Appendix Table A1).

The main result. Passed measures are followed by issuance 35.6% of the time against 14.6% for failed ones, but most of that 21-point gap is selection: better-supported projects both pass and proceed. The design removes the selection. At the cutoff, authorisation raises six-year GO issuance by **+11.0 percentage points (RBC, robust z = 2.30; conventional +14.6, z = 4.47)**; any-issuance gives +11.6 and +14.4 (Table 8, Figure 3). The effect replicates state by state

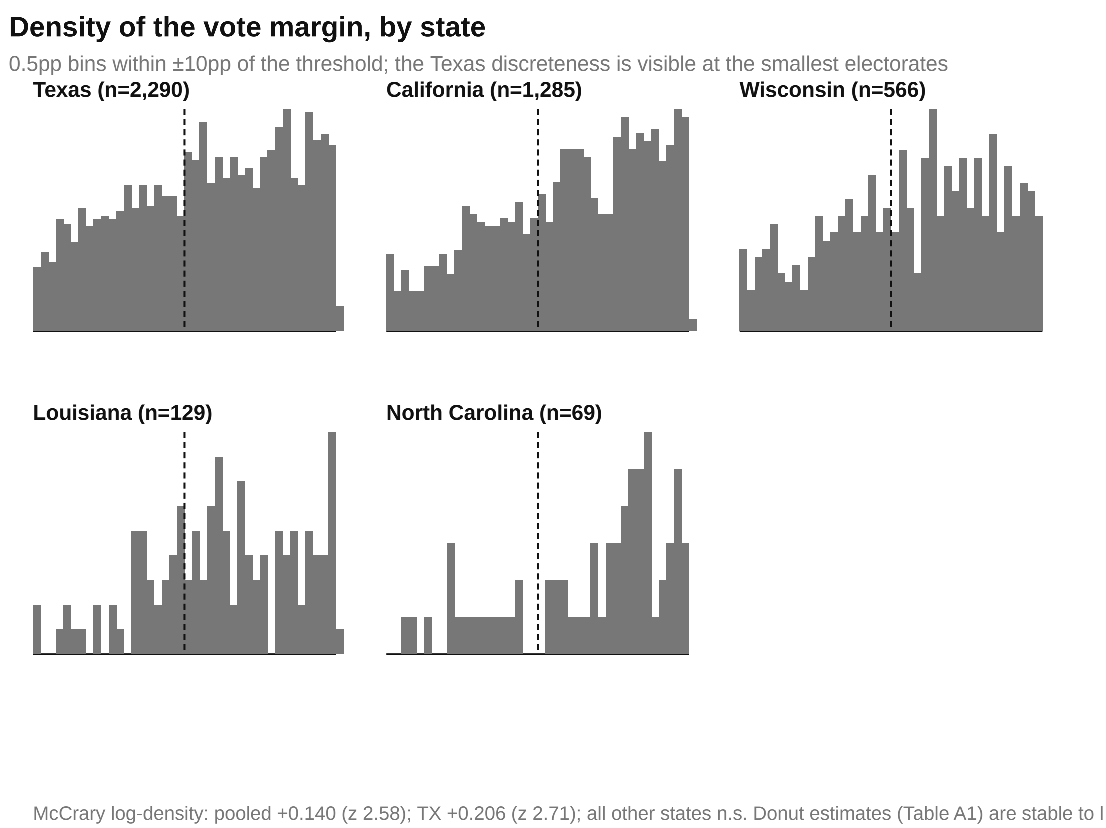

Figure 2: Density of the vote margin by state. Histograms of the running variable within ±10pp; the discreteness of small Texas electorates is visible. McCrary log-density discontinuity: pooled +0.24 (z 4.6) at h = 5pp, +0.14 (z 2.6) over this full ±10pp window; Texas +0.21 (z 2.7); all other states approximately zero.

across three different thresholds: California +16.6 (z 3.0), Texas +14.1 (z 3.3), Wisconsin +21.2 (z 3.1). Three statutes, one answer. Louisiana's negative cell (−27.6, z −1.9) is a data-grain artefact, measures folding to the parish, and excluding it moves the pooled estimate to +16.2 (Appendix Table A2 reports the state cells; Appendix Table A3 the placebo thresholds, all null at the 5% level).

Table 8: The effect of authorisation on borrowing and building

Outcome	RBC	robust p	Conv.	h	Eff. N (L/R)
Any GO issuance ≤6y	+0.110** [+0.016, +0.204]	0.022	+0.146	±10	1,596/2,743
Any issuance ≤6y	+0.116** [+0.023, +0.209]	0.014	+0.144	±10	1,596/2,743
ln(1+new-money par p.c.)	+0.937** [+0.057, +1.818]	0.037	+0.918	±10	1,537/2,627
ln(1+LTD issued p.c.), Census survey	+0.878** [+0.146, +1.610]	0.019	+0.827	±10	1,238/2,128
Same-purpose continuation ≤6y	+0.052 [-0.047, +0.152]	0.302	+0.073	±10	1,398/2,456
Voter-mode first issuance ≤6y	+0.184*** [+0.089, +0.278]	<0.001	+0.199	±10	1,596/2,743
No issuance ≤6y	-0.116** [-0.209, -0.023]	0.014	-0.144	±10	1,596/2,743
Capital outlay p.c. (pay-go bound; schools, post−pre)	+0.464** [+0.090, +0.837]	0.015	+0.377	±10	996/1,621

Notes: Frame: rd_sample ∩ GO measures, corpus package v3 (frozen). RBC = robust bias-corrected local-quadratic estimate at h with robust CI (Calonico–Cattaneo–Titiunik, $\rho = 1$); conventional = local-linear, triangular kernel. Bandwidth fixed at ±10pp per the reconciliation in Appendix Table A1 (MSE-optimal row and bandwidth-sensitivity curve there). * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$ (robust). Survey row: GFD/IUF, EMMA-independent (survivorship check). Continuation: ballot purpose matched to use-of-proceeds functions; bridge precision 80.0% (blind audit). Pay-go row: full-window cohort (votes to ∼2017–18, outlay to FY2023).

Reading. A narrow authorisation raises six-year GO issuance by eleven points (robust), roughly doubles per-capita borrowing in disclosure and survey data alike, raises same-purpose project delivery, and moves building nearly one-for-one with borrowing: refusal is not offset by pay-go construction.

The intensive margin doubles, in two independent sources. Log new-money par per capita rises +0.92 (z 3.0) in the disclosure corpus and, decisively for the survivorship concern, **+0.83 (z 3.3) in the Census survey measure**, which includes the bank loans and private placements no official statement records. Both are RBC-stable (+0.94, +0.88) and the survey pre-period placebo is null (+0.41, z 1.5). This is the survey debt items' single appearance, as Section 4's architecture prescribes.

Timing, not denial. The event study (Figure 4) loads the entire effect on the vote year (τ_0 = +0.236, z = 7.8; every pre- and post-year approximately zero). Figure 5 turns the timing into quantities for the |margin| ≤ 5 window: among six-year issuers, the median barely-authorised government reaches the market in **0.33 years**, the median barely-refused one in **1.15**, a delay of 0.8 years, and the refused side never reaches the authorised side's end-of-vote-year issuance level within six years. By year six, 50.1% of barely-refused governments have still issued nothing, against 40.0% of barely-authorised ones: a ten-point wedge that is the durable residue after all catch-up.

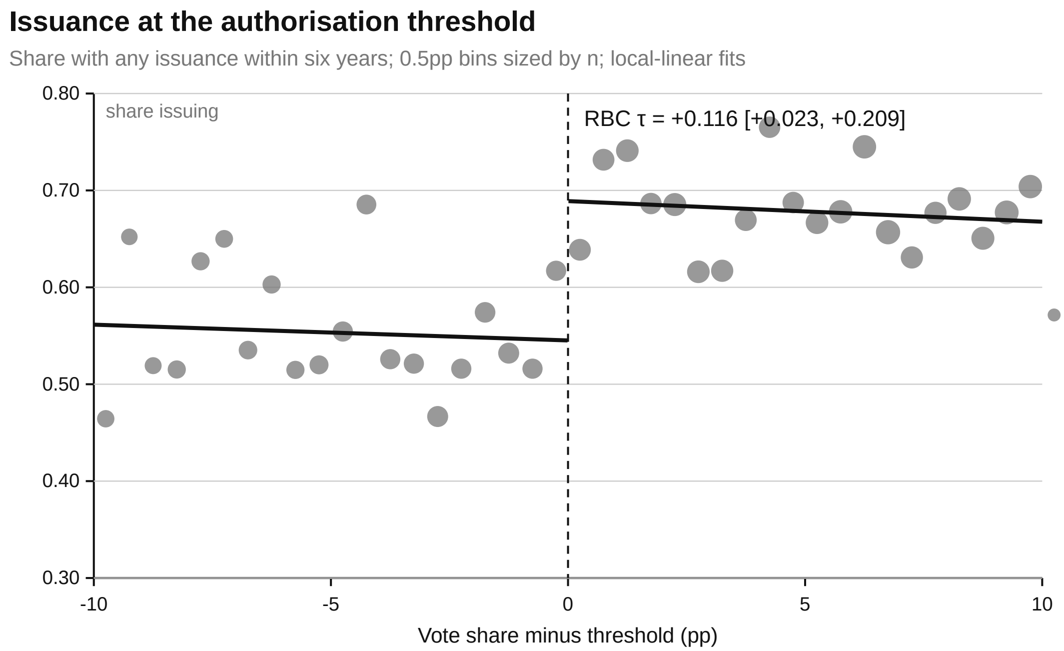

Figure 3: Issuance at the authorisation threshold. Share of measures followed by any issuance within six years, in 0.5pp bins of the running variable sized by cell count, with local-linear fits on each side and the RBC estimate annotated.

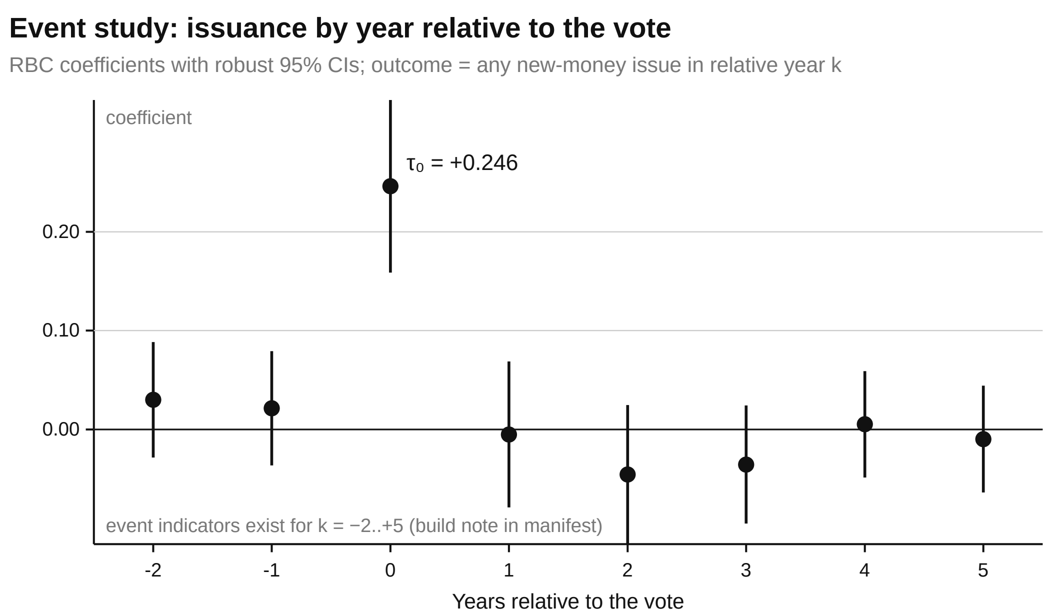

Figure 4: Event study. RBC coefficient on issuance in each year relative to the vote, with robust 95% confidence intervals; indicators exist for k = −2 to +5.

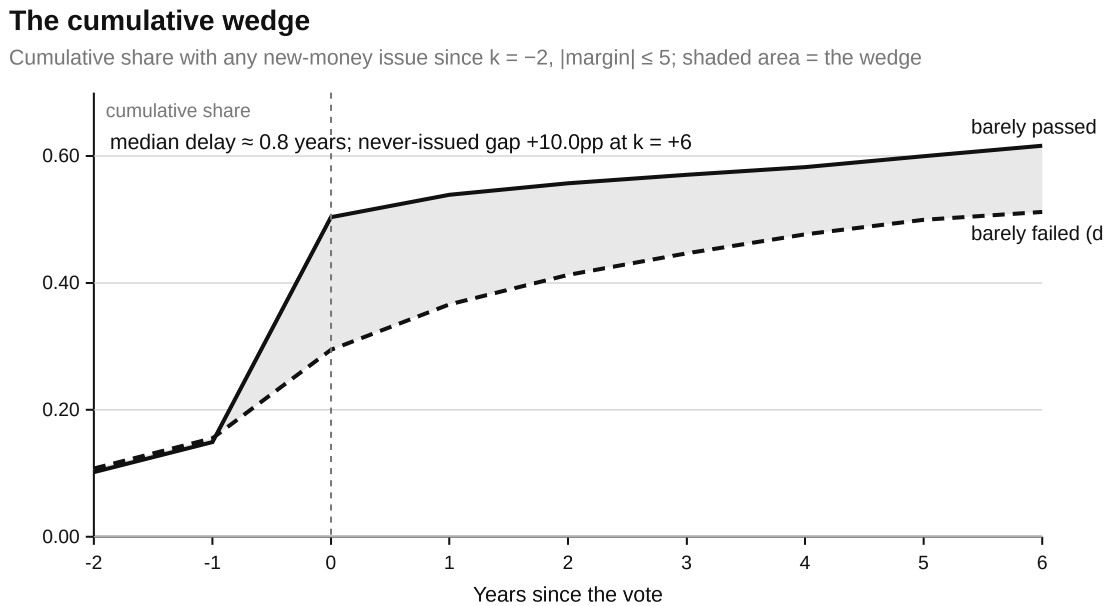

Figure 5: The cumulative wedge. Cumulative any-issuance for barely-passed (solid) and barely-failed (dashed) measures, |margin| $\leq$ 5; the shaded area is the wedge; median market-entry times annotated.

Robustness, and the pay-go bound. The full specification battery is Appendix Table A1: donut variants, clustering by unit (z 4.11) and county (z 3.47), Lee bounds for crosswalk selection at trim 2.52% ([+0.138, +0.163]), randomisation inference in the ± 2 window ($p < 0.0002$; 0 of 5,000 permutations), horizons one to six years, and the bandwidth curve. Two diagnostics do not reject and are reported with their reconciliation: the IK bandwidth selects h = 1.8pp, only two points of support, where slope terms absorb a level shift that the design-based tests at the same window detect decisively; RBC at ± 5 has the same character. The battery's final row is the **pay-go bound**: rerunning the design with survey capital outlay per capita as the outcome, the schools post-minus-pre differenced estimate is **+0.377 (z 2.9; RBC +0.464, z 2.4)** beside an issuance effect of +1.43, and trimmed levels imply an outlay-to-issuance ratio of **0.92** at the cutoff. Refused governments do not detectably replace lost borrowing with debt-free construction. (Caveats: outlay spreads over construction years, so six-year windows understate long-project differences; the differenced specification nets the positive pre-period outlay RD that endogenous timing produces; total outlay only; the full-window cohort is votes through roughly 2017 to 2018 with outlay ending at fiscal 2023.) Nothing in the battery is hidden; the appendix table shows all of it.

7 · The response margin: what refusal buys

The theory's distinctive claim is that a coalition requirement prices delay, not denial. This section follows every refused measure forward and asks what refusal actually bought. Table 9 collects the accounting.

Refusal is a pause. Of 2,680 failed GO measures, the re-submission hazard is front-loaded:

Table 9: The response margin: what refusal buys

Quantity	Value
Panel A: the fate of 100 marginal refusals (2005–19 cohort, $\|m\| \leq 5$, n=422)	
Re-approved by voters ≤4y	54.3
Returned, not yet converted	13.3
Issued via board or statutory channel	5.2
Issued on pre-existing voter authority	9.0
Extinguished within horizon	18.2
First post-vote authoriser ($\|m\| \leq 5$): barely-passed vs barely-failed	
barely-PASSED: voter 48.4%, board 7.9%, none 43.1%	
barely-FAILED: voter 36.7%, board 8.6%, none 54.2%	
Panel B: re-submission (2,680 failed GO measures)	
Hazard, years 1/2/3/4 (%)	26.7 / 22.8 / 15.8 / 12.4
Cumulative return within 4y (%)	58.2
Median time to return (years)	1.02
Returns that pass (%)	61.9
Median amount ratio, return/original	1.000
Purpose categories retained on return (%)	78.2

Notes: Fates are mutually exclusive per 100 barely-refused measures; issuance on pre-existing voter authority draws on authorisations banked before the refused measure and is not evidence of evading it. Panel B: full failed sample; amounts where registries print them (n=1,354). Retention from the audited purpose bridge (precision 80.0%).

Reading. Refusal is a pause: most marginal refusals are re-approved at the next election at the same ask, fewer than one in five extinguish, and the board channel is a floor rather than the treatment margin.

26.7% return within a year, then 22.8%, 15.8% and 12.4% in the following years, cumulating to **58.2% within four years**, with the median return at **1.02 years**, which is to say the next election. Returns win: **61.9%** pass. And districts re-ask rather than concede: the median returning measure asks for 100% of the original amount (n = 1,354; only 45.9% downsize).

The fate of the marginal refusal. For the 2005 to 2019 cohort with full observation windows, per 100 barely-refused measures: **54 are re-approved by voters within four years; 13 more return and await a verdict; 5.2 borrow through the board or statutory channel; 9.0 issue on pre-existing voter authority; 18 are extinguished** within the horizon. The 9.0 issuing on pre-existing authority draw on authorisations banked before the refused measure, so they are not evidence of evading it; the substitution reading attaches only to the 5.2 board-or-statutory cell. The transition matrix beneath Table 9 shows what the cutoff actually moves: voter-mode first issuance (48.4% for barely-passed against 36.7% for barely-failed) against no issuance (43.1% against 54.2%), while board-mode first issuance sits near 8% on both sides. The board channel is a floor, not the treatment margin. Consistent with this, what barely-refused governments still issue tilts towards the board channel (council share $\tau = -0.064$, z = -2.0).

The project itself survives. Matching ballot purposes to the use-of-proceeds functions of subsequent issues (a deterministic category bridge, blind-audited at **80.0% precision and 88.9% recall**), same-purpose financing appears within six years for 44.7% of barely-passed measures and **33.4% of barely-failed ones** (RD +0.072, z 2.1; imprecise under RBC: +0.052, robust p = 0.30), with median arrival 0.32 against 1.67 years. A third of narrowly refused projects get financed anyway, later. On re-submission, bundles are not recomposed: 78.2% of original purpose categories are retained. The pay-go bound of Section 6 closes the remaining loophole on the provision side:

building falls with borrowing at a ratio near one, so the continuation gap is not being quietly filled by debt-free construction. [**PENDING: nine disputed audit pairs queued for review.**]

Authorisation banked, drawdown deferred. The supermajority regime produces a distinct fate for near-misses. Within California, failures with majority support short of the threshold re-submit far more than decisive failures (42.2% against 27.4%; 51.3% against 30.8% among schools) and show no jump at the symbolic 50% line: the response tracks proximity to passage, not the majority label. Yet their issuance is lower, and the deficit is not a truncation artefact: it **widens** from −5.3 points at six years to −6.4 at eight (schools −6.7 to −8.0). The near-miss chain (Appendix Table A8) shows why. Near-miss failures convert by re-vote four times as often (107 against 25 passed returns), but the conversions sit undrawn: median pass-to-first-issue **6.2 years** against 2.9, and only 5.5% see a voter-mode issue within eight years. Where the bar is a supermajority, even re-assembled consent does not become borrowing on the study horizon. A candidate mechanism, that Proposition 39's statutory tax-rate caps meter the drawdown of 55%-route authorisations, was tested and is neither supported nor excluded: the uncapped cell is too thin (9 conversions against 118) and its point direction runs against the cap story (Appendix Table A9). The fact stands documented and unexplained.

8 · Where the requirement binds: exits and electorates

The RD average of Section 6 conceals the theory's structure. Two partitions recover it: by the borrowing government's exit menu, and by the electorate that the rule convenes.

By exit menu: the fork. Table 10 splits the RD frame by entity class and sets each class's estimate against its independently measured national menu from Section 5. School districts (non-voted share 29.8% within the frame) show the binding effect (+0.147, z 4.0; RBC +0.123, z 2.3) and the highest re-submission rate (59.2%). General-purpose governments (non-voted share 80.4%) show no discontinuity at all (+0.073, z 0.9; RBC −0.02) and the lowest re-submission (25.3%). Special districts sit between, with the largest GO-specific jump, which shrinks at any-issuance as part of the gap reroutes to non-GO instruments; their cells are small and the ordering, not the magnitude, is the cited object. The rule binds exactly where the menu offers no exit; where exits exist, refusal is nearly costless and rarely re-litigated. This is the same fork the fifty-state substitution results (Table 6) show in cross-section.

By electorate: the stable propertied public. A naive assembly-cost reading predicts that the authorisation effect weakens wherever coalitions are harder to build, and in particular that demographic heterogeneity raises the price of assembly. The data falsify that version on its own sign: the effect is carried by homogeneous electorates (+0.232 against −0.018 in diverse ones), not blocked by them. What the moderator profile consistently favours is a **stable-propertied-public** account: authorisation binds where the consenting public is durable and propertied, because those are the electorates whose refusals stay refused and whose approvals licence long-horizon drawdown. Homeownership above the frame median gives τ = **+0.182 (z 3.5)** against +0.056 below (Table 11); the 65-plus split gives +0.194 against +0.036 in levels (the formal difference is not significant under RBC); the effect is stronger on-cycle (+0.252 against +0.099), when the broad durable electorate is the one consulted. It is ownership, not affluence: the income split is flat (+0.129 against +0.122). And it is institutional, not partisan: precinct-built city partisanship and mayoral party in the 577-city panel show no moderation (all n.s.; the coarser county-grain splits agree and sit in Appendix Table A12). All moderation tables run on the 6,255 proper-

Table 10: Exits and the binding of refusal: the fork against the menu

Outcome	RBC	robust p	Conv.	h	Eff. N (L/R)
Schools, GO issuance	$+0.123^{**}$ [+0.018, +0.228]	0.022	+0.147	±10	1,289/2,144
Any issuance	$+0.132^{**}$ [+0.027, +0.236]	0.013	+0.150	±10	1,289/2,144
menu: non-voted $ share 36.6%; failures re-submitting 0.592					
General-purpose, GO issuance	-0.024 [-0.251, +0.203]	0.837	+0.073	±10	252/498
Any issuance	-0.028 [-0.249, +0.192]	0.800	+0.045	±10	252/498
menu: non-voted $ share 86.1%; failures re-submitting 0.253					
Utilities (special districts), GO issuance	$+0.500^{*}$ [-0.015, +1.016]	0.057	+0.463	±10	55/101
Any issuance	+0.359 [-0.169, +0.888]	0.182	+0.362	±10	55/101
menu: non-voted $ share 73.8%; failures re-submitting 0.426					

Notes: Frame: rd_sample ∩ GO measures, corpus package v3 (frozen). RBC = robust bias-corrected local-quadratic estimate at h with robust CI (Calonico–Cattaneo–Titiunik, $\rho = 1$); conventional = local-linear, triangular kernel. Bandwidth fixed at ±10pp per the reconciliation in Appendix Table A1 (MSE-optimal row and bandwidth-sensitivity curve there). * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$ (robust). Menu shares from the national corpus (independent of the RD frame). Utilities cells are small; the ordering, not the magnitude, is the cited object.

Reading. Refusal binds exactly where the exit menu is poorest: schools show the discontinuity and the highest re-submission; general-purpose governments, with the richest menu, show none.

grain measures (place grain for cities, district grain for schools); special districts are excluded pending sub-county demographics. The reframing is stated as it is: an account adopted after seeing the homogeneity sign, offered with its transmission story, and flagged as post hoc rather than pre-registered.

Table 11: Where the requirement binds: splits at within-frame medians

Outcome	RBC	robust p	Conv.	h	Eff. N (L/R)
Homeownership: below median	-0.065 [-0.232, +0.103]	0.451	+0.056	±10	432/863
Homeownership: above median	+0.196** [+0.044, +0.348]	0.011	+0.182	±10	575/861
difference (above−below)	+0.261**	0.024			
Share 65+: below median	+0.043 [-0.129, +0.215]	0.622	+0.036	±10	408/781
Share 65+: above median	+0.097 [-0.055, +0.249]	0.210	+0.194	±10	599/943
difference (above−below)	+0.054	0.646			
Racial-ethnic homogeneity (1−frac.): below median	-0.092 [-0.270, +0.086]	0.311	-0.018	±10	415/818
Racial-ethnic homogeneity (1−frac.): above median	+0.203*** [+0.059, +0.348]	0.006	+0.232	±10	592/906
difference (above−below)	+0.295**	0.012			
Median household income: below median	+0.030 [-0.127, +0.187]	0.711	+0.129	±10	482/802
Median household income: above median	+0.123 [-0.041, +0.288]	0.142	+0.122	±10	525/922
difference (above−below)	+0.094	0.420			
On-cycle (November, even year)	+0.192** [+0.024, +0.361]	0.025	+0.252	±10	476/1,005
Off-cycle	+0.072 [-0.039, +0.182]	0.205	+0.099	±10	1,120/1,738
difference (on−off)	+0.121	0.240			

Notes: Frame: rd_sample ∩ GO measures, corpus package v3 (frozen). RBC = robust bias-corrected local-quadratic estimate at h with robust CI (Calonico–Cattaneo–Titiunik, $\rho = 1$); conventional = local-linear, triangular kernel. Bandwidth fixed at ±10pp per the reconciliation in Appendix Table A1 (MSE-optimal row and bandwidth-sensitivity curve there). * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$ (robust). Demographic splits run on the 6,255 proper-grain measures (place grain for municipalities, district grain for schools); special districts excluded pending sub-county demographics. Partisanship (precinct-built 577-city panel, all subgroups n.s.) in Appendix Table A-P1.

Reading. The effect concentrates where the consenting public is propertied, older, homogeneous and consulted on-cycle; it is ownership, not income, and institutional, not partisan.

The incidence of the higher bar. Blocked majorities, measures with majority support short of a supermajority threshold (California only, by construction), arise in places that sit between decisive-failure and comfortable-passage places on renter share, diversity and inequality; among schools they are the poorer places (median household income −\$3.8k against cleared places, SE 1.5k; Appendix Table A10). Descriptive, but it locates the supermajority's demographic cost:

majorities in less affluent, mid-composition districts.

9 · The agenda margin and the politics of the rule

Rules discipline what is asked. The coalition requirement operates before any vote is cast. Under California's 55% bar, school districts bring **8.8 proposals per 100 districts per year** at a median $39M, 62.7% on-cycle; under Texas's 50% bar, **20.1** per 100 at $15M, 14.1% on-cycle; Wisconsin sits between on all three margins (Table 12). Fewer, larger, better timed. Pass rates make the same point in reverse: nearly invariant between the 50% and 55% regimes (79.2% against 77.6%) and collapsing only at two-thirds (47.6%). A moderate bar is absorbed by the agenda; an extreme bar defeats even what selection permits. The Texas 2019 reform (HB 3) closes the loop within-state: when the state mandated separate propositions for stadiums, natatoria and performing-arts facilities, propositions per election jumped from 1.5 to 1.7 before the reform to 2.0 to 2.4 after, and the multi-proposition share doubled within two cycles. Bundling is a rule-governed choice, not a habit. The near-cutoff excess mass of Section 6 belongs to the same family: proposals are timed to win.

Table 12: The agenda margin by regime

state (thr)	#districts	proposals/yr	per 100 districts/yr	median amount	on-cycle share
CA (55)	1065	93.5	8.78	$39,000,000	62.7%
TX (50)	1093	219.2	20.05	$15,000,000	14.1%
WI (50)	440	48.4	11.00	$12,600,000	30.5%
MN (50)	392	0	–	–	–

Notes: School GO proposal behaviour 2005–25; MN empty by classification (a data fact). Passage-rate and TX-2019 unbundling panels in the CSV companion (A4b, A4c).

Reading. Under the higher bar districts propose half as often, at 2.6 times the size, four times more on-cycle; pass rates barely move until two-thirds.

At fixed voter support, the rule alone moves outcomes. In the 50 to 55% band, the identical electoral result is a failure for a Californian school district and a success for a Texan or Wisconsin one. The difference-in-differences across the adjacent band gives **+11.5 points of six-year issuance (SE 4.8)** attributable to the rule at fixed support, indistinguishable from the RD estimate obtained from an entirely different comparison: an out-of-design calibration regarded here as the strongest single check in the paper.

The polity fights over the rule exactly where the theory predicts. Since 1990, nearly every attempt to lower a local borrowing threshold has targeted schools, the class Table 2 shows holding the poorest exit menu: California's Propositions 170 (1993, failed), 26 (2000, failed) and 39 (2000, passed 53 to 47, schools only, with a $30M elite-financed bipartisan campaign); Washington's 2023 to 2025 school-bond bills (all stalled); Idaho's roughly eleven legislative attempts (none passed). The one recent non-school attempt, California's Proposition 5 (2024), which would have extended the 55% bar to housing and infrastructure, failed. Where exits exist, no one spends thirty million dollars to lower the bar. [**PENDING: the reform table is compiled from secondary sources and marked secondary-unverified pending primary-record checks; it is cited as evidence on reform politics, not as a causal estimate.**]

Appendix tables and figures

Table A1: Specification battery for the main estimate

Check	Result
Donut $\|m\| > 0.5$pp	RBC +0.208 [+0.087, +0.329]; conv. +0.195; N 1,502/2,622
Donut $\|m\| > 1.0$pp	RBC +0.188 [+0.033, +0.343]; conv. +0.186; N 1,409/2,498
Donut $\|m\| > 2.0$pp	RBC +0.249 [-0.024, +0.521]; conv. +0.199; N 1,199/2,240
Excluding Louisiana	RBC +0.133 [+0.037, +0.229]; conv. +0.162
Cluster by unit (2,480)	conv. +0.146, SE 0.036, z 4.11
Cluster by county (650)	conv. +0.146, SE 0.042, z 3.47
Lee bounds (crosswalk trim 2.52%)	[+0.138, +0.163]
Randomisation inference, $\|m\| \leq 2$	diff. +0.142, $p < 0.0002$ (5,000 permutations)
IK MSE-optimal bandwidth	$h = 1.81$pp; $\tau = -0.001$ (SE 0.075): slope-dominated on two points of support; the design-based tests above at the same window reject decisively
Alternative horizons 1–6y	Appendix Figure A1a (coefficients rise to the 6y wedge)
Bandwidth sensitivity $h \in [3, 15]$	Appendix Figure A1b (stable from $h \geq 6$)

Notes: Frame: rd_sample ∩ GO measures, corpus package v3 (frozen). RBC = robust bias-corrected local-quadratic estimate at h with robust CI (Calonico–Cattaneo–Titiunik, $\rho = 1$); conventional = local-linear, triangular kernel. Bandwidth fixed at ±10pp per the reconciliation in Appendix Table A1 (MSE-optimal row and bandwidth-sensitivity curve there). * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$ (robust). Clustered, Lee, randomisation and IK rows from the committed polish battery (POLISH_RESULTS.md).

Reading. The estimate survives every design-based and inference variant; the one non-rejecting diagnostic (the MSE-optimal local fit) is a power statement, reconciled in the note.

Table A2: The main estimate, state by state

Outcome	RBC	robust p	Conv.	h	Eff. N (L/R)
Texas (50%)	$+0.134^{**}$ [+0.011, +0.256]	0.032	+0.141	±10	892/1,398
California schools (55%)	+0.059 [-0.154, +0.271]	0.589	+0.105	±10	261/692
California non-school (66.7%) (small cell)	$+0.295^{**}$ [+0.048, +0.541]	0.019	+0.260	±10	166/164
Wisconsin (50%)	$+0.167^{*}$ [-0.021, +0.356]	0.082	+0.212	±10	221/345
Louisiana (50%) (small; parish-fold grain)	-0.396^{*} [-0.796, +0.004]	0.052	-0.276	±10	40/89
North Carolina (50%) (small cell)	-0.209 [-1.353, +0.934]	0.720	+0.216	±10	16/53

Notes: Frame: rd_sample ∩ GO measures, corpus package v3 (frozen). RBC = robust bias-corrected local-quadratic estimate at h with robust CI (Calonico–Cattaneo–Titiunik, $\rho = 1$); conventional = local-linear, triangular kernel. Bandwidth fixed at ±10pp per the reconciliation in Appendix Table A1 (MSE-optimal row and bandwidth-sensitivity curve there). * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$ (robust). Louisiana outcomes fold to the parish; the pooled estimate excluding Louisiana appears in Appendix Table A1.

Reading. Three adequately powered states, three different thresholds, one answer; the small cells are directionally consistent except fold-grain Louisiana.

Table A3: Placebo thresholds: no jump where the law draws no line

Outcome	RBC	robust p	Conv.	h	Eff. N (L/R)
50%-regime states, pseudo-cutoff 45 (both sides fail)	$+0.143^{*}$ [-0.004, +0.290]	0.057	+0.012	±10	786/687
50%-regime states, pseudo-cutoff 60 (both sides pass)	-0.045 [-0.129, +0.039]	0.296	-0.014	±10	1,874/1,852
CA 66.7% regime, pseudo-cutoff 60 (both sides fail)	-0.068 [-0.394, +0.259]	0.685	-0.007	±10	69/126
CA 55% regime, pseudo-cutoff 60 (both sides pass)	-0.039 [-0.264, +0.187]	0.736	+0.052	±10	280/764

Notes: Frame: rd_sample ∩ GO measures, corpus package v3 (frozen). RBC = robust bias-corrected local-quadratic estimate at h with robust CI (Calonico–Cattaneo–Titiunik, $\rho = 1$); conventional = local-linear, triangular kernel. Bandwidth fixed at ±10pp per the reconciliation in Appendix Table A1 (MSE-optimal row and bandwidth-sensitivity curve there). * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$ (robust). Pseudo-cutoffs estimated within samples where both sides share the same legal outcome.

Reading. Issuance is smooth through vote shares at which nothing legal changes; the discontinuity appears only at the statutory threshold.

Table A4: Classified-line coverage by regime (nc-share selection check)

sample	regime	units	unit coverage	$ coverage
all classes	strict	12,838	63.5%	82.7%
all classes	non-strict	2,040	77.6%	83.9%
all classes	not codable	3,450	48.3%	80.1%
general-purpose (muni+county)	strict	3,776	68.7%	90.3%
general-purpose (muni+county)	non-strict	1,006	75.1%	84.9%
general-purpose (muni+county)	not codable	1,926	60.9%	92.5%
schools	strict	6,763	66.7%	80.7%
schools	non-strict	263	90.9%	96.1%
schools	not codable	604	3.8%	7.7%

Notes: Coverage of B3-classified project dollars among corpus-active units; regime labels PRELIMINARY. Count-based composition versions accompany every dollar-based exhibit.

Reading. Coverage is regime-unbalanced on the unit margin, so the text cites the count basis; both bases agree on direction.

Table A5: Chargeable share of classified project dollars by authorisation channel

auth mode	$ch (B)	$nc (B)	chargeable share
voter	33.8	264.1	11.3%
council_or_board	658.6	439.9	60.0%
statutory	183.8	69.7	72.5%
refunding_no_new_election	73.5	47.4	60.8%
unknown	72.0	34.0	67.9%

Notes: National corpus, classified printed-amount use lines; count-based version in Table A-C1b.

Reading. The voted channel finances the non-chargeable civic core; the unvoted channels carry the chargeable perimeter.

Table A6: Observed voted share by rule regime and entity class

entity	rule	units	voted $ share	council $ share	GO security share	non-chargeable share
county	strict	2,375	23.3%	70.9%	49.1%	61.6%
county	non-strict	481	5.4%	90.7%	80.7%	63.1%
municipal	strict	11,037	22.2%	74.1%	37.7%	34.2%
municipal	non-strict	2,305	3.7%	89.2%	72.1%	52.6%
township	strict	5,735	18.0%	80.2%	89.3%	61.0%
township	non-strict	2,483	46.5%	51.8%	97.1%	75.9%
special_district	strict	23,373	33.3%	57.1%	33.0%	30.7%
special_district	non-strict	6,801	13.1%	83.8%	33.4%	18.2%
school_district	strict	12,162	68.4%	30.6%	85.4%	99.4%
school_district	non-strict	616	7.4%	90.6%	45.2%	99.8%

Notes: Dollar-weighted shares from the entity panel (v3); rules PRELIMINARY; townships carry a proxy municipality rule (the reversal is a coding artefact scheduled for pass-2).

Reading. Raw magnitudes behind Table 4: under strict rules two-thirds of school dollars are voted; under lax rules almost none are.

Table A7: Institutional and demographic interactions with the rule (appendix)

	(1) Voted \$ share TEL interaction	(2) GO share TEL interaction	(3) NC share TEL interaction	(4) Voted \$ share × homeownership	(5) Voted \$ share × county Dem share
Strict rule (main effect)	+0.205	-0.109	+0.005	+0.096	+0.194
Interaction	-0.185	-1.227	-0.638	-0.047	-0.651**
	(0.338)	(0.821)	(0.494)	(0.395)	(0.326)
Observations	396	401	358	13,928	13,928
State clusters	34	34	34	48	48

Notes: Columns (1)–(3): the roughly 570-city subpanel with tax-and-expenditure-limit (TEL) stringency; the interaction asks whether TELs sharpen or blunt the rule; the subpanel is the exit-rich class where weak rule effects are predicted, so power is limited by design. Columns (4)–(5): national entity panel, strict rule interacted with the centred moderator; these moderate the first-stage channel, not the RD issuance effect (the causal moderator tests are in Section 8). Entity panel (90,604 local governments; estimation samples below are the corpus-active units for each outcome). Weighted least squares; Census-region fixed effects; controls: ln population (ln total revenue for special districts), homeownership, share 65+, racial fractionalisation, ln median household income, county Democratic presidential share 2020. Standard errors clustered by state in parentheses. The strict-rule indicator is the PRELIMINARY pass-1 coding (verification pass in progress), so coefficients read as first-stage associations, not causal estimates. * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.

Table A8: The near-miss chain

years since failure k	maj-fail issued ≤k (n at risk)	min-fail issued ≤k (n)	gap
1	1.9% (n=371)	4.6% (n=174)	-2.7%
2	4.9% (n=327)	8.4% (n=155)	-3.5%
3	6.7% (n=326)	10.3% (n=155)	-3.6%
4	7.9% (n=303)	12.8% (n=148)	-4.9%
5	10.6% (n=303)	15.0% (n=147)	-4.4%
6	10.7% (n=270)	16.0% (n=106)	-5.3%
7	13.0% (n=270)	18.9% (n=106)	-5.9%
8	14.5% (n=256)	20.8% (n=96)	-6.4%

Notes: CA failed GO measures; per-cell observability restrictions; new-money issues only.

Reading. The near-miss deficit widens with the window: not a truncation artefact.

Table A9: The rate-cap split (inconclusive)

split	group	conversions	median pass→first issue (y, n)	issued ≤6y of pass (n)	issued ≤8y of failure (n)
A – within schools (route choice – selection flagged)					
	schools, 55% return (rate-capped)	118	4.93 (64)	35.6% (104)	34.6% (104)
	schools, 66.7% return (uncapped)	9	9.74 (8)	22.2% (9)	33.3% (9)
B – cross-class (regime $\not\perp$ entity – CONFOUNDED, flagged)					
	schools (capped 55)	118	4.93 (64)	35.6% (104)	34.6% (104)
	non-schools (uncapped 66.7)	5	12.59 (2)	25.0% (4)	25.0% (4)

Notes: 132 conversions; the uncapped cell (14; 9 school) is too thin to test and points against the cap story. Verdict: mechanism neither supported nor excluded.

Reading. Re-approved but unissued stands as a documented fact without an adjudicated mechanism.

Table A10: Demography of blocked majorities (descriptive; California)

covariate	blocked	cleared	sub-50	blocked-cleared (SE)	blocked-sub50 (SE)
renter share	0.370 (n=277)	0.408 (n=1526)	0.351 (n=161)	-0.038 (0.006)	+0.019 (0.010)
Gini (B19083)	0.442 (n=277)	0.450 (n=1526)	0.437 (n=161)	-0.008 (0.002)	+0.005 (0.004)
fractionalization	0.538 (n=277)	0.558 (n=1526)	0.511 (n=161)	-0.020 (0.010)	+0.027 (0.014)
median HH income $	64,710 (n=277)	68,508 (n=1526)	68,344 (n=161)	-3,798 (1,463)	-3,635 (2,284)
65+ share	0.130 (n=277)	0.130 (n=1526)	0.142 (n=161)	+0.000 (0.003)	-0.012 (0.005)
SAIPE child poverty (schools)	0.155 (n=107)	0.175 (n=696)	0.142 (n=95)	-0.020 (0.010)	+0.013 (0.012)

Notes: Blocked = majority short of the supermajority threshold; within-matched ACS/SAIPE at the D5 grain ladder; DESCRIPTIVE, no causal claim.

Reading. Blocked majorities arise in less affluent, mid-composition school districts: the incidence of the higher bar.

Table A11: Validation of extracted authorisation fields

state	T1: docs w/ date	exact match	±7d	matched→PASSED	T2: voter docs	supported ≤6y	T3: council docs	no recent passed GO
CA	779	62.8%	64.1%	98.0%	849	66.4%	1238	91.0%
LA	417	62.8%	63.3%	86.6%	412	86.7%	208	79.8%
MA	260	34.6%	40.0%	94.4%	1619	76.9%	756	100.0%
MN	209	19.6%	19.6%	85.4%	247	22.7%	1362	100.0%
NC	258	65.9%	65.9%	98.8%	0	–	0	–
TX	114	73.7%	73.7%	94.0%	7671	76.8%	4454	78.3%
WI	900	95.4%	95.6%	87.1%	821	90.5%	443	86.0%
pooled	2937	67.9%	68.8%	91.3%	11619	76.2%	8461	86.1%

Notes: Election-date matches against independently observed referenda; support and consistency tests as defined in the text. GFD bridge: 2022 public-use issuance matches GFD within 0.5% for 99.9% of bridged units.

Reading. Where registries are complete the extracted fields track the election record at ninety per cent or better; shortfalls sit exactly where registry coverage is known-short.

Table A12: County-grain partisanship (demoted from the text)

subgroup	n	τ GO-issue	SE	z
Dem share < median (0.423)	5073	+0.121	0.047	2.60
Dem share $\geq$ median	5118	+0.111	0.048	2.30
tercile 1 (most Republican)	3397	+0.163	0.055	2.94
tercile 2	3397	+0.072	0.058	1.23
tercile 3 (most Democratic)	3397	+0.120	0.060	2.01

Notes: County presidential two-party share, 2020 vintage, county grain: a coarse proxy for district electorates. The text's partisanship null cites the precinct-built 577-city panel. National first-stage interaction: strict $\times$ county Dem -0.65 (t -2.0), descriptive.

Reading. The county-grain splits agree with the city-panel null; nothing partisan moderates the authorisation effect.

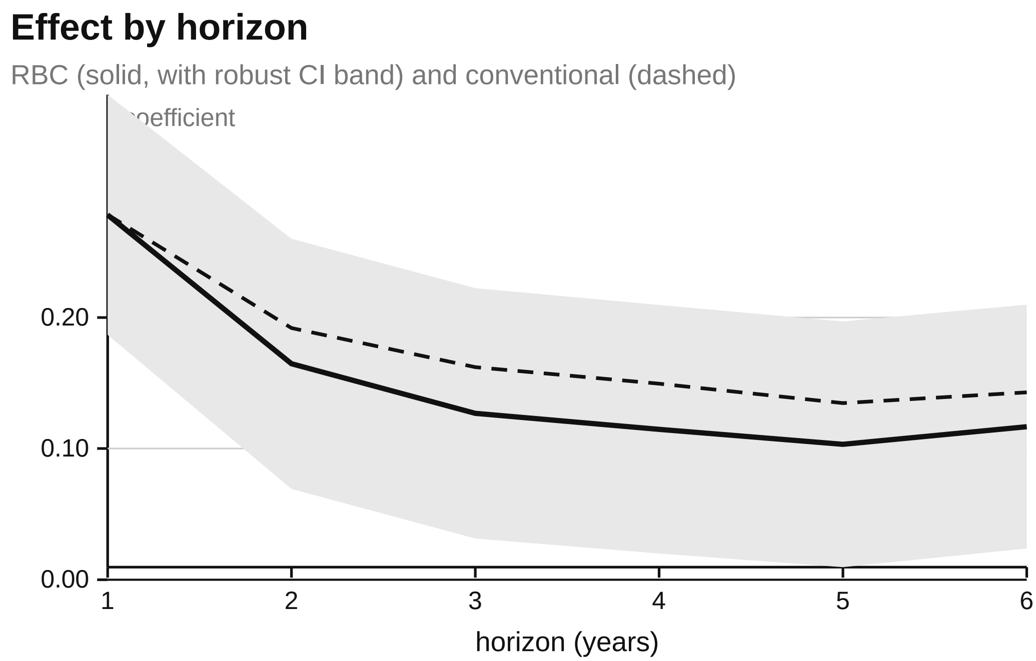

Figure A1: Effect by issuance horizon: RBC estimate and robust CI for windows of one to six years.

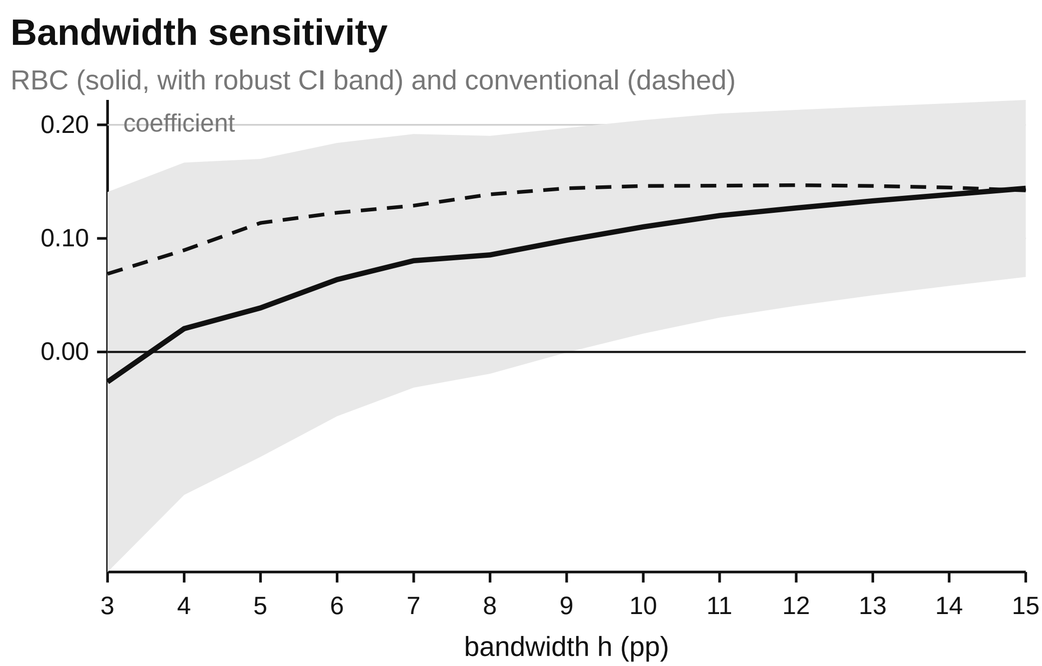

Figure A2: Bandwidth sensitivity: RBC estimate (solid), conventional (dashed) and robust CI across bandwidths h of 3 to 15pp.

Appendix H · The honesty record

Moved verbatim from the main text per the round-3 adjudication, updated for P4/P5. The register of informative nulls, pending items and known data limits that the paper carries.

Informative nulls

- Partisanship (three measures, two grains: county presidential, precinct city-footprint, mayoral party): no moderation of the authorisation effect.
- The extensive margin nationally: the coded rule does not predict whether a government borrows at all, only the channel and composition.
- Entity midwifery: no detectable spawning of new issuers in refused units' counties (coarse grain, low power).
- The D2b signal test: no jump in the response at the symbolic 50% line among institutional failures.
- TEL by rule on the big-city subpanel: null, power-limited by design (large cities are the exit-rich class).
- P4 rate-cap split: the banked-authorisation mechanism is neither supported nor excluded (9 uncapped conversions; point direction opposite); the fact stands unexplained.

Pending items, each blocking a specific claim

- Rules human pass-2 (21 cells, with the owner): C2/H2 finals; causal upgrades of the fifty-state first-stage and substitution results; the township rule column (town-meeting states break the municipality proxy).
- Proposition 39 district-level design (state matching-fund controls; CDIAC pre-2005 issuance for the composition margin): journal-version item, not to be started before the conference draft ships.
- B5 audit disagreement review (9 pairs; upgrades the 80% bridge precision).
- Ten-year-plus horizon pass on the re-approved-but-unissued cohort.
- Special-district demographics beyond county proxy (TWDB and shapefile interpolation).
- R1 reform table: secondary_unverified pending the owner's primary-record verification.
- IL and IN vote margins (harvest scoped and paused).

Appendix tables added in round 4

- **Table A-C1 (nc_share selection check):** classified-amount-line coverage by state and regime, unit- and dollar-weighted (`NC_COVERAGE_RESULTS.md`). Coverage is regime-UNBALANCED on the unit margin (schools 24.1pp, general-purpose 6.4pp), so every composition exhibit reports a count-based version beside the dollar-based one; the text cites the count-based general-purpose coefficient (-0.047, t -2.25) and both channel-sorting gradients (dollars 11.3/60.0/72.5; counts 17.9/34.5/54.0).

- **Table A-P1 (county partisanship, demoted from the text):** county presidential two-party splits of the authorisation effect (+0.121 against +0.111; terciles +0.163/+0.072/+0.120, non-monotone) and the national first-stage interaction (-0.65, t -2.0). Caveat: 2020 vintage only, county grain (a coarse proxy for district electorates). The text's partisanship null cites the precinct-built 577-city panel only.
- **Moderation-table note (all D5/H3 tables):** estimates run on the 6,255 proper-grain measures (place for cities, SAIPE/ACS district grain for schools); special districts excluded pending sub-county demographics (TWDB interpolation).
- **Fate-table labels:** "issued anyway" is split into board/statutory channel (5.2) and pre-existing voter authority (9.0); the latter is not evidence of evasion of the refused measure.
- **Version pin:** the conference draft is frozen on corpus package v3; package v4 and the B5 rerun are journal-revision items.
- **Pay-go bound cohort:** the outlay row's full-window cohort is votes through roughly 2017–18 with outlay ending at fiscal 2023.

Known data limits, stated where used

- Louisiana parish-fold grain (outcome mixes measures within a parish).
- Minnesota school-purpose classification gap (empty E4 cell is a classification fact, not evidence).
- GFD security split (FFC/NG) unreported after 2005 for every entity type and degenerate for schools before (FFC is 100% of school long-term debt by classification); the corpus security class carries that outcome.
- GFD capital outlay: total only (construction not itemised in the compact; F-33 not in the repo); no outlay items in the 2024 public-use continuation, so outlay windows end at fiscal 2023.
- Texas BRB counts-unknown placeholder rows ("1-0", 3,188 rows) excluded from coalition-size panels; no RD estimate affected.
- Corpus truncation at 2005 (EMMA era).
- Moderator coverage is within-matched (ACS and SAIPE bridges); matched subsets are not representative of unmatched units.
- Consent-map coverage: ten states carry no determined-mode dollars (an extraction-coverage fact, shown as missing rather than imputed).

Appendix V · Variable inventory by section (§§4–9)

Reference document, 2026-08-24, computed on the current panels (corpus package v3). "Panel" = the 40,924 crosswalked referenda; "RD frame" = rd_sample ∩ bond_go, 11,889 measures. Coverage figures are non-blank counts from the committed files, not estimates. Conditional variables state their own denominator. British spelling; no em dashes.

§4 · Data (identity, treatment definition, validation)

variable	definition	level	source	coverage	notes
`unit_id` / GID	9-char Census government ID (state2+type1+county3+unit3)	unit	crosswalk	40,924/47,235 referenda (86.6%)	exact tiers ~100% verified; fuzzy tier 95.1% RA-audited
`pct_yes`	yes share of the vote	referendum	9 state registries	75% of panel; **100% of RD frame**	IL/IN carry results only (no margin)
`threshold_centered_margin`	pct_yes − applicable threshold	referendum	derived; CA threshold per measure from CDIAC	75% / **100%**	the running variable
`rd_sample`	at a genuine mandatory-ballot statutory cutoff	referendum	rules coding of the 9 states (direct, not the national panel)	23,577 flagged	RD-state cutoffs are bedrock statutes, outside the pass-2 worklist
`purpose_class`	bond_go / bond_other / tax etc.	referendum	per-state purpose mapping	100%	MN school classification gap flagged
`auth_mode_final2`	OS-evidenced authorisation mode (voter / council_or_board / statutory)	document	corpus extraction	93.7% of docs determined (v3)	validated: date match 67.9% pooled, 95.4% WI; matched→passed 91.3%
`op_referendum_strict` (+ ordinal, threshold_num)	coded rule, state × entity × purpose × year	state×entity×year	rules panel	89% of panel rows; muni cell codable in 37 states	**PRELIMINARY pass-1; 21-cell pass-2 with owner**

§5 · Landscape and fifty-state first stage

Grain: one row per local government (entity panel, 90,604 units) or per document (corpus).

variable	definition	level	source	coverage	notes
`nm_docs`, `nm_par`	new-money issues and par, 2005–25, canonical per issue	unit	corpus v3	78,672 typed nm issues; 46 states ≥50 docs	5 small states below gate (DC/DE/HI/VT/WY); 5 legacy states 72–87% flag coverage (MN/MA/MO/MD/ID)
`sec_go_sh` ... `sec_sptax_sh`	par shares by security pledge class	unit	corpus	where nm_par>0	GO / revenue / lease / special_tax

variable	definition	level	source	coverage	notes
`voted_sh_par`, `council_sh_par`, `statutory_sh_par`	mode shares of determined nm $	unit	corpus	13,875 units with a determined split	the first-stage outcome
`nc_share_project`	non-chargeable share of classified project $	unit	B3 doc flags (118-label map, 0 unmapped)	~18% of units	classified printed-amount lines only
ballot purposes (M2 panel A)	normalised keyword categories	referendum	9 registries	19,600 bond measures	stand-in for a formal B1 taxonomy (flagged)
votes cast / yes-needed (M3)	counts and threshold×total	referendum	TX/WI/LA/NC registries	7,473 GO measures	CA has % only; 3,188 TX placeholder rows excluded; LA also carries voters_qualified/voted
GFD characteristics	pop or enrolment, revenue, own-source, property tax, LTD out	unit, latest yr ≥2012	national GFD	100% of entity panel	specials lack Population (revenue-size proxy in N-suite)
ACS covariates	homeownership, 65+, fractionalisation, median income	place (munis) / county (rest)	ACS5-2019 national pulls	~100%	county layer uniform; place upgrade for cities
`county_dem2p_2020`	county Dem two-party share	county	MEDSL	99–100%	2020 only
FOG institutions	form of gov, initiative, referendum, partisan, districts	city	FOG panel	~570 big cities	subpanel only
TEL stringency; ACS-2022; city Dem; mayor party	big-city extras	city	municipal-analysis panels	~570 cities each	2013-vintage TEL, no time variation
rule (per class)	strict for the unit's entity class	state×entity	rules panel	county 94% · schools 93% · specials 79% · munis 68% · townships 50% (proxy)	township column = pass-2 Part C

§6 · The RD core (frame: 11,889 GO measures)

variable	definition	level	source	coverage in RD frame	window
`issued_6y`	any issuance ≤6y	referendum	corpus link	**100%**	(0, +6y]
GO issuance ≤6y	issued_6y ∧ go_share>0	referendum	corpus	100% (go_share observed for 64%, else 0-issue)	headline outcome
`nm_par_6y`, `ln_par_pc_6y`	new-money par; ln(1+par p.c.)	referendum	corpus + GFD denominator	100% / 86%	denominator = gfd_pop, enrolment for schools (65% of frame are schools)
`ln_gfd_ltd_pc_6y`	survey LTD issued p.c. (EMMA-independent row)	referendum	GFD + IUF FY23/24	73%	fiscal [y+1,y+6]; IUF validated 99.9%

variable	definition	level	source	coverage in RD frame	window
capital outlay p.c. (pay-go bound row)	GFD Total_Capital_Outlays p.c.	referendum	GFD	83% with a 6y window (8,489/10,244 with denominator)	ends FY2023 (IUF has no outlay); total outlay only
`ev_m2...ev_p5`	any nm issue in relative year k	referendum×year	corpus	100%	event study
balance covariates	pre-vote GFD pop, revenue p.c., debt p.c., taxes, enrolment, prior failure	referendum	GFD	94–98%	0/7 imbalanced
McCrary/density inputs	binned margin counts	bin	derived	full frame	TX-driven excess

§7 · The response margin (denominators are conditional)

variable	definition	level	source	coverage	notes
`resubmitted_4y`, hazard, KM	same-unit+class return ≤4y	failed measure	panel sequences	2,680 failed GO measures (the denominator)	risk sets censored at registry end
return pass / amount ratio	outcome of the return	return	registries	1,418 returns; 1,354 with amounts	amounts CA/TX/WI
fate table categories	converted / returned / issued-anyway / extinguished	failed measure	corpus + sequences	1,233 (2005–19 full windows); 422 close	mutually exclusive
first-issue authoriser	mode of first post-vote nm issue	referendum	corpus	close window n=2,234	transition matrix
continuation (B5)	window doc sharing the ballot's purpose category	measure×doc	FN2CAT bridge	3,854 close categorisable measures	bridge precision 80.0% / recall 88.9% (blind audit); reruns on v4
bundle recomposition	categories kept/dropped on return	failure→return pair	bridge	1,129 pairs	
P3 chain	pass→first-issue timing; 8y issuance	conversion	corpus	132 CA conversions (118 capped / 14 uncapped)	uncapped cell too thin (P4 not confirmed)

§8 · Heterogeneity (moderators; split at within-frame medians)

variable	definition	level (grain)	source	coverage in RD frame	notes
`acs_homeown/share65/frac/medinc`	proper-grain moderators	**place 2,103 · SD 4,152 · county 5,329**	ACS5 2010/2019 (pre-vote vintage rule)	97% attached	proper-grain-only splits use the 6,255 place+SD rows
`sd_childpov_rate`	district child poverty (income proxy)	school district	SAIPE name-bridge	35% of frame (4,153 schools)	match rates vary (MN 91% . . . CA 44%); within-matched validity
county moderators	65+, frac, income (first-pass proxy)	county	CC-EST + SAIPE	97.4%	superseded by proper grain where available
`county_dem2p` (pre-vote)	partisanship	county	MEDSL	85.7%	null result

variable	definition	level (grain)	source	coverage in RD frame	notes
city Dem share, mayor party	large-city partisanship	city (precinct-built)	municipal-analysis	1,932 municipal measures in 577-city panel; 178 within ±10	replicate-the-null only
Gini (D6)	B19083	place/SD/county	new CA pull (2010+2019)	CA cells (blocked-majority analysis)	D6 is CA-only by construction
election timing	on-cycle = Nov of even year	referendum	dates	100%	
entity class (fork)	census_type	unit	crosswalk	100%	schools 7,690 · GP 2,623 · specials 1,576

§9 · Agenda margin and politics

variable	definition	level	source	coverage	notes
proposals /100 districts/yr; median ask; on-cycle share	E4 regime comparison	state×year	registries + GFD district counts	CA/TX/WI (MN = classification gap)	
pass rates by regime	E3	regime	registries	11,888 with margins	invariance at 50 vs 55; breaks at 66.7
props/election; multi-prop share	TX-2019 unbundling	district×election	TX BRB	2014–25	reporting-mandate caveat
W1/W2 wedge cells	issuance by support band × regime	band×state group	panel	50–55 band: CA 187 / TX+WI 653	the DiD-at-fixed-support
reform events	attempts/outcomes since 1990	state×event	web compilation	11 rows	secondary_unverified (owner)
Prop 39 / M56 DiD outcomes	ln per-pupil (per-capita) LTD issued	state×year cells	national GFD 1997–2013	schools ≥20 districts/cell, 39–48 states	state-mean grain; district-level upgrade deferred

The three load-bearing coverage facts a referee will probe

1. **RD outcomes are 100%/86% covered in-frame** (binary/intensive) and the survey replication row (73%) is symmetric across the cutoff; crosswalk selection is Lee-bounded. 2. **Moderator grains are honest**: place for cities, district for schools (SAIPE/ACS-SD), county fallback (all specials); splits are within-matched and labelled; nothing is imputed. 3. **National corpus coverage post-v3**: 46 states above the 50-doc gate; the residuals are five thin small states and five legacy partial-fill states, both listed; rule coding 68–94% by class with the township proxy flagged.